

Midterm Exam

- ◉ Date: **10:45 am – 11:45 am, Oct 31, 2023, Tuesday**
- ◉ Classroom: CYCP1
- ◉ Format:
 - ◉ Open-book, but no computers or any electronics
 - ◉ About 5 questions (some are short questions)
 - ◉ Answer sheets will be provided
 - ◉ Calculators are allowed, but optional
- ◉ Content:
 - ◉ Everything so far cut at “Deadlock” (excluded)
- ◉ Absence Policy:
 - ◉ 10% marks automatically merged to Final Exam
- ◉ Solutions to PS#2 will be released in Monday midnight.

Deadlock

2023-2024 Fall COMP3230A

Contents

- ◉ What is deadlock?
- ◉ Four necessary conditions of deadlock
- ◉ Deadlock prevention
- ◉ Deadlock avoidance
- ◉ Detection and recovery

Related Learning Outcome

- ◉ ILO 2c - explain the underlying causes of deadlock issues and describe the principles and techniques used by OS to support concurrency control

Readings & References

- ◉ Required Reading
 - ◉ Chapter 32 – **Common Concurrency Problems**
 - ◉ <http://pages.cs.wisc.edu/~remzi/OSTEP/threads-bugs.pdf>

Producer/Consumer Problem: Problem?

```
sem_t empty, full, mutex;
```

```
sem_init(&empty, 0, MAX); // MAX buffers are empty initially
```

```
sem_init(&full, 0, 0); // 0 buffers are full
```

```
sem_init(&mutex, 0, 1); // mutex = 1 because it is a lock (binary semaphore)
```

```
Producer()
```


```
{  
    while (1)  
    {  
        P(&mutex); /* lock */  
        P(&empty); /* block if buffer unavail */  
  
        // critical section  
        <<< Put item into shared buffer >>>  
  
        V(&full);  
        V(&mutex);  
    }  
}
```

```
Consumer()
```

```
{  
    while (1)  
    {  
        P(&mutex);  
        P(&full);  
  
        // critical section  
        <<< Remove item from shared buffer >>>  
  
        V(&empty);  
        V(&mutex);  
    }  
}
```

Deadlock Problem in Real Life


It's a system-wide tangle of **resource requests**, but the system is in a state that **all requests cannot be fulfilled**, and **the whole system** comes in a **standstill**.



The Deadlock Problem

- ⦿ A set of threads **each holding** some system resources and **block waiting** to acquire another system resource **held by another thread** in the set.
- ⦿ Examples

<i>T0</i>	<i>T1</i>
<i>wait (A);</i>	<i>wait(B)</i>
<i>wait (B);</i>	<i>wait(A)</i>



<i>T0</i>	<i>T1</i>
<i>lock (&lock1);</i>	<i>lock (&lock1);</i>
<i>lock(&lock2);</i>	<i>lock(&lock2);</i>
<i>while (wait) {</i>	<i>signal(&cv, &lock2);</i>
<i>wait(&cv, &lock2);}</i>	<i>unlock(&lock2);</i>
<i>unlock(&lock2);</i>	<i>unlock(&lock1);</i>
<i>unlock(&lock1);</i>	

<i>T0</i>	<i>T1</i>
<i>buffer1.put();</i>	<i>buffer2.put();</i>
<i>buffer1.put();</i>	<i>buffer2.put();</i>
<i>buffer2.get();</i>	<i>buffer1.get();</i>
<i>buffer2.get();</i>	<i>buffer1.get();</i>

Why do Deadlocks occur?

- ◉ In previous examples, if we detect the situations, probably deadlock would not happen
- ◉ Unfortunately, sometimes **we cannot see how threads use their locks**
 - ◉ in large code bases with complex dependencies
 - ◉ in programs with **external library functions**, as with encapsulation that hides the details of implementation

Vector vec1, vec2;

Thread1
:
:
vec1.AddAll(vec2);
:


Thread2
:
:
vec2.AddAll(vec1);
:
:

To avoid race condition,
AddAll() probably needs to
lock both Vectors before
performing the computation.

Classical Synchronization Problem

- ◉ Dining Philosophers


- ◉ Five philosophers sit around a circular table. Each leads a simple life alternating between thinking and eating spaghetti. In front of each philosopher is a dish of spaghetti that is constantly replenished by a dedicated wait staff. There are exactly five forks on the table, one between each adjacent pair of philosophers. Eating spaghetti (in the most proper manner) requires that a philosopher uses both adjacent forks (simultaneously). Develop a concurrent program free of deadlock and indefinite postponement that models the activities of the philosophers.



Dining Philosophers

“The problem is famous because it is fun and somewhat intellectually interesting; however, its practical utility is low.”

The key challenge is to show that your solution is **without deadlock**, **no philosopher is being starved**, and **concurrency is high**



philosopher

```
while (true) {  
    think();  
    eat();  
}
```

```
eat() {  
    pickupLeftFork();  
    pickupRightFork();  
    eatingSpaghetti();  
    putdownRightFork();  
    putdownLeftFork();  
}
```




Is this solution working?

Necessary Conditions for Deadlock

- ◉ Four conditions must be held for a deadlock to occur; if any of these are **not met**, deadlock **CANNOT** occur
 - ◉ **Mutual exclusion condition**
 - ◉ Allow only one thread to have exclusive access to a resource
 - ◉ Wait-for condition (**hold-and-wait condition**)
 - ◉ A thread may hold some resources while awaiting assignment of additional resources
 - ◉ **No-preemption condition**
 - ◉ No resource can be forcibly removed from a thread that holding it
 - ◉ **Circular-wait condition**
 - ◉ Two or more threads are locked in a “circular chain” in which each thread is waiting for one or more resources that the next thread in the chain is holding
- ◉ The four conditions are **necessary but not sufficient**

Circular-Wait Condition

- Thread T1 has been allocated resource R4 that is being requested by thread T2 that has been allocated resource R3 that is being requested by T1





Deadlock vs. Starvation

- ◉ **Starvation:** A thread fails to make progress for an indefinite period
- ◉ **Deadlock:** a group of threads forms a cycle where none make progress because each is waiting for some others to act
- ◉ **Deadlock → Starvation; Starvation !→ Deadlock**
 - ◉ E.g., Readers-Writers problem: Starving/**waiting** writers are waiting for **active** readers to finish
 - ◉ **Any active thread waiting on one or more waiting threads?**
- ◉ A thread **subject to starvation/deadlock** does not always have starvation/deadlock

Deadlock Prevention

- By using **restrictive policy** in allocation of resources to **remove any one** of the four necessary conditions, **deadlock cannot happen**
- Prevent Circular-wait condition
 - Imposes a total ordering** of all resource types, and **requires** that each threads **requests resources in** an increasing **order** of enumeration
 - Disadvantage:
 - Not all programs using resources in that order, but you are required to hold resources of smaller labels before granting resources of larger labels; this may lead to poor resource utilization



Deadlock Prevention

- ◉ Prevent Hold-and-wait condition
 - ◉ At start, thread gets all needed resources all at once or nothing
 - ◉ So the hold-and-wait condition is never satisfied
 - ◉ Disadvantages:
 - ◉ Low resource utilization
 - ◉ Starvation possible
 - ◉ a process requests many resources may have to wait for a longer time as this strategy favors waiting processes with small resource needs
- ◉ The Mutual Exclusion condition
 - ◉ Sharable resources, if allow non-mutually exclusive access, do not result in deadlock
 - ◉ Unfortunately, most sharable resources don't support or work properly under non-mutual exclusive access

Deadlock Prevention

- ◉ Denying No-preemption condition
 - ◉ If a process that is holding some resources requests another resource but not immediately available, must **release all holding resources**
 - ◉ Process will be restarted only when it can regain its old resources, as well as the new ones that it is requesting
 - ◉ Disadvantages:
 - ◉ This can lead to substantial overhead
 - ◉ when a process releases resources, it may lose all of its work to that point; or it may have to undo all previous work before going to restart
 - ◉ possibility of starvation


Deadlock Avoidance

- ◉ Does not preventively remove one of the conditions for deadlock; instead, the system **tries to avoid deadlock** if it **knows ahead of time** **all the resources requests** associated **with each** of the threads
- ◉ Avoidance **requires some global knowledge** of which locks/resources various threads might grab during their execution
- ◉ Subsequently schedules the threads in a way as to guarantee no deadlock can occur
- ◉ Two approaches
 - ◉ Avoidance by scheduling
 - ◉ Banker's algorithm

Avoidance by Scheduling

- Given the resources needs of different concurrent threads, the system **looks at their dependency** and use it as a guideline in scheduling the threads to avoid deadlock
- Example

	Thr1	Thr2	Thr3	Thr4
R1	need	need	no	no
R2	need	need	need	no



- If the system allows Thr1 & Thr2 to execute concurrently by different CPUs, there is a possibility of having deadlock
- This approach is a bit **conservative** and may result in **under utilization** of the resources

Banker's Algorithm



- ◉ Threads are **allowed to hold locks/resources** while **requesting** additional locks/resources
- ◉ System only grants the allocation of additional resources to a thread when the allocation **will not** result in an **unsafe state**
 - ◉ i.e., the system estimates that deadlock would not happen **with the remaining resources**
 - ◉ by means of checking whether with the remaining resources, there still **exists a feasible allocation solution** to satisfy **all future demands from all threads** that could lead to successful termination of all threads
- ◉ It has a number of weaknesses, such as
 - ◉ **requiring to know ahead of time the resource needs** as well as only works with a fixed number of threads and resources
 - ◉ e.g., if a device breaks and not available, the algorithm won't work as this may cause the state to turn to unsafe

Banker's Algorithm

Declared
maximum
usage

Current
usage

Future
requests

◉ Safe, unsafe, deadlocked states

	max()	loan()	claim()
T1	4	1	3
T2	6	4	2
T3	8	5	3
total=12		available=2	

Total unit of
resources

Remain
available

If T1 asks for 1 unit, should the system grant this request?

If T2 asks for 1 unit, should the system grant this request?

**In this example, we only
have one type of resource**

Detection & Recovery

- ◉ Allow deadlocks to occasionally occur
 - ◉ “If a bad thing happens rarely, certainly one should not spend a great deal of effort to prevent it, particularly if the cost of the bad thing occurring is small”
- ◉ System takes action **periodically** to check whether deadlock has happened
 - ◉ Identifies processes and resources involved in the deadlock
 - ◉ Usually focus on determining if a circular wait exists
 - ◉ One technique for detecting deadlocks involves building a **resource-allocation graph** and looking for **cycles** (necessary but not sufficient if some resources have multiple instances)
- ◉ Recovery
 - ◉ Solution 0: Process without the resource
 - ◉ Solution 1: Abort all deadlocked processes
 - ◉ Solution 2: Abort one process at a time until the deadlock cycle is eliminated.
 - ◉ Successively preempt resources until deadlock no longer exists
 - ◉ Transactions: rollback and retry
 - ◉ No matter what, some processes become victims
 - ◉ Removal generally requires that the process be restarted from beginning or from a previous checkpoint

Summary

- ◉ Deadlock is a serious issue commonly found in concurrent programs, such as OS, multithreaded programs, and highly parallel programs.
- ◉ When we identify a deadlock scenario, we always find the four necessary conditions appeared in that scenario
- ◉ Solutions to deadlock
 - ◉ Prevention – use restrictive rules or guidelines to deny one of the necessary conditions
 - ◉ it is the responsibility of the programmers to apply the rules / guidelines in their programs
 - ◉ Avoidance – require to have the global knowledge of locks / resources usage amongst all threads, and use these as the hint to allocate locks / resources to threads
 - ◉ Programmers are free to arrange their logic in resource acquisition; it is the responsibility of the systems to apply the strategy in avoiding deadlock
 - ◉ Detection and Recovery – a pragmatic solution with no much overhead most of the time, but need more effort in recovering from deadlock

Non-Deadlock Concurrency Issues

- ◉ Atomicity violation
 - ◉ “The desired serializability among multiple memory accesses is violated.”
- ◉ Order violation
 - ◉ “The desired order between two (groups of) memory accesses is flipped.”

Application	What it does	Non-Deadlock	Deadlock
MySQL	Database Server	14	9
Apache	Web Server	13	4
Mozilla	Web Browser	41	16
OpenOffice	Office Suite	6	2
Total		74	31

Figure 32.1: **Bugs In Modern Applications**

Operating Systems

◉ Virtualization

- ◉ CPU Virtualization
 - ◉ Process Abstract
 - ◉ Address space
 - ◉ Process states
 - ◉ Process control block
 - ◉ Process operations API
 - ◉ Signals
 - ◉ Limited Direct Execution
 - ◉ System calls
 - ◉ Context switch
 - ◉ Interrupts
 - ◉ Scheduling
 - ◉ Scheduling metrics
 - ◉ FIFO, SJF, HRRN, STCF, RR, MLFQ
 - ◉ Multi-core scheduling, Linux CFS
- ◉ Memory Virtualization
 - ◉ Address space
 - ◉ Address translation: dynamic relocation
 - ◉ Segmentation
 - ◉ Paging
 - ◉ TLB
 - ◉ Multi-level paging
 - ◉ Inverted page table
 - ◉ Swap space
 - ◉ Page replacement policy: FIFO, LFR, LRU, Clock
 - ◉ Thrashing

◉ Concurrency

- ◉ Thread
 - ◉ POSIX threads (pthreads)
 - ◉ Race conditions, critical sections, mutual exclusion, atomic operations, synchronization
- ◉ Locks
 - ◉ Atomic instructions: test-and-set, compare-and-swap
 - ◉ Mutex locks
- ◉ Condition Variables
 - ◉ Pthread CVs
 - ◉ Producer-Consumer problem
- ◉ Semaphores
 - ◉ Binary Semaphores
 - ◉ Counting Semaphores
 - ◉ Ordering
 - ◉ Readers-Writers problem
- ◉ Deadlock
 - ◉ Dining philosophers' problem
 - ◉ Four necessary conditions
 - ◉ Deadlock prevention, avoidance, detection&recovery

◉ Persistence

- ◉ I/O devices (HDD, SSD)
- ◉ Files and Directories
 - ◉ Inode
 - ◉ File descriptor
 - ◉ Hard/Symbolic links
- ◉ File System Implementation
 - ◉ On-disk data structure
 - ◉ Superblock, Bitmap, Inodes, Data blocks
 - ◉ Free space management
 - ◉ Bitmap, linked-list, block-list
 - ◉ Caching and buffering
 - ◉ Access control and protection
 - ◉ Journaling file system
 - ◉ Data journaling
 - ◉ Metadata journaling
- ◉ **Advanced Topics**