

# **Mechanism – Paging: Finding Free Page Frame**

2023 Fall COMP3230A

# Not Enough Physical Memory

- ◉ Assumption 1
  - ◉ Assume the address space of a process is relatively small, i.e., size of physical memory is much larger than a process's address space.
- ◉ In real life, process's address space is quite large (e.g., 32-bit  $\rightarrow$  4 GiB), but physical memory is limited
- ◉ The Crux
  - ◉ If the system does not have enough physical memory, how to run many processes at the same time?

# Contents

- ◉ To support **multiple** running processes, each with large address space, OS uses storage disks to **temporarily** store **portions** of processes' address spaces
  - ◉ Swap Space and Page Fault
- ◉ We need to move out some virtual pages and make way for newly request or access pages
  - ◉ Replacement Policies
  - ◉ Evaluation of the policies
- ◉ Thrashing

# Related Learning Outcomes

- ◉ ILO 2b - describe the principles and **techniques** used by OS in **effectively** virtualizing memory resources.
- ◉ ILO 3 [Performance] - **analyze and evaluate** the algorithms of . . . and **explain** the major performance issues . . .

# Readings & References

- ◉ Required Readings
  - ◉ Chapter 21 – **Beyond Physical Memory: Mechanisms**
    - ◉ <http://pages.cs.wisc.edu/~remzi/OSTEP/vm-beyondphys.pdf>
  - ◉ Chapter 22 – **Beyond Physical Memory: Policies**
    - ◉ <http://pages.cs.wisc.edu/~remzi/OSTEP/vm-beyondphys-policy.pdf>

# The Crux

- ◉ How can OS make use of a larger, slower device to *transparently* provide the illusion of a **large virtual address space**?

# Swap Space

- ◉ Most OSs create a special area (partition) of the disk as **swap space**
  - ◉ OS **swaps** virtual pages **out** of physical memory to it and **swaps** virtual pages **back** into physical memory from it
- ◉ To mitigate the performance overhead due to swapping
  - ◉ This partition is not associated with/managed by file management
  - ◉ Consists of **consecutive tracks** to increase disk read/write performance

# Memory Hierarchy


- Tradeoff: speed, size, and cost of storage
  - Smaller, faster; slower, cheaper
- An example:

Cache	Hit Cost	Size
1 <sup>st</sup> level cache / 1 <sup>st</sup> level TLB	1 ns	64 KB
2 <sup>nd</sup> level cache / 2 <sup>nd</sup> level TLB	4 ns	256 KB
3 <sup>rd</sup> level cache	12 ns	2 MB
Memory	100 ns	10 GB
Data center memory (the same LAN)	100 us	100 TB
Local non-volatile memory	100 us	100 GB
Local disk	10 ms	1 TB
Data center disk	10 ms	100 PB
Remote data center disk	200 ms	1 XB



# Page Fault

- ◉ Present bit



- ◉ Indicates whether this virtual page is **in physical memory or on disk**
- ◉ During address translation,
  - ◉ if the processor finds that this page is a **valid** page, but the page's **present bit is zero**, means not in physical memory, the processor generates a **page fault** (a type of exception)
  - ◉ That triggers the OS to invoke the **page-fault handler**, to load the missing page from secondary storage into physical memory
- ◉ When the page is not in memory, OS needs to know the disk address of the page in the swap space
  - ◉ Usually, the disk address is **stored in the PTE**; probably share with the bits used for storing the PFN

# Page Fault

- ◉ OS (page-fault handler) finds the disk address in the PTE and **issues an I/O request** to fetch the page into memory
- ◉ While the I/O is in flight, OS **places** the current process in **blocked state**, and selects another ready process to run
- ◉ When I/O completes, OS
  - ◉ **Updates** the PFN field and present bit of the corresponding **PTE**
  - ◉ Then **unblock the process** and **triggers** the process to **retry the instruction** (which triggers this page fault)
    - ◉ This causes a TLB miss, and TLB hardware fetches the new translation information from the PTE and updates the TLB cache
    - ◉ Then retry the instruction again

# Finding Free Page Frames

- ◉ OS page-fault handler needs to find a free page frame for placing the incoming page. Where to find it?
- ◉ Get it from the Free-list
  - ◉ As page is a fixed-size block and same size of a frame, any free frame in anywhere should be okay
- ◉ What if memory is full?
  - ◉ OS has to first **swap out one or more virtual pages** to make room for the requested page(s)

# Page Replacement Policy

- ◉ Strategy used by OS to decides **which virtual page to move out** from main memory to make space for incoming page
- ◉ The replacement policy is **critical** to the performance of the application
  - ◉ If select a wrong page to kick out, this will result in experiencing more page faults
  - ◉ Thus, using inappropriate policy can cause a program to run at disk-like speeds, which is **much much much** slower than the CPU speed

# Replacement Policy

- ◉ Replacement policy is being used in various system components, and the goal is to **minimize** the number of **misses** or to **maximize** the **hits**
  - ◉ In cache management → cache miss
  - ◉ In TLB management → TLB miss
  - ◉ In virtual memory management → page fault
- ◉ One common way to assess the effectiveness of policy is to measure the **hit rate** and calculate the average memory access time (**AMAT**), which is defined as  $(Hit_{\%} \cdot T_M) + (Miss_{\%} \cdot T_D)$ 
  - ◉ Where  $T_M$  is the cost of accessing memory when hit, and  $T_D$  is the cost of accessing memory when missed (including get back the page from disk and access the memory in that page)

# Find a “victim” page

- ⦿ A replacement strategy is characterized by
  - ⦿ The **heuristic** it uses to select a page for replacement
  - ⦿ **Execution overhead** it incurs

# Evaluation

- ◉ Evaluate individual policy by running the algorithm on a particular sequence of memory references (called **reference string**) and computing the **hit rate** of that reference string
- ◉ In all our examples, the reference string is
  - ◉ 2 3 2 1 5 2 4 5 3 2 5 2
- ◉ When assessing individual policy, we assume a process has **3 page frames**, and all are **initially empty**

# Optimal Replacement

- ◉ The victim page is the page that **will be** accessed **furthest in the future** (as compare to other in memory pages)
  - ◉ i.e., will not be referenced for longest period of time in the future
- ◉ It **always** leads to the **fewest misses** overall
  - ◉ **Impossible to have perfect knowledge of future events**
  - ◉ Acts as a baseline for comparing how well a policy performs

2	3	2	1	5	2	4	5	3	2	5	2
2	2	2	2	2	2	4	4	4	4	4	4
	3	3	3	3	3	3	3	3	2	2	2
			1	5	5	5	5	5	5	5	5
F	F		F	F		F			F		

Cold-start misses

- ◉ There are 6 faults with this reference string; thus, the **hit rate is 50%**



# First-In-First-Out Replacement

- ◉ The victim page is the page that has been in the system the longest
  - ◉ pages are placed in a queue and **the oldest page** is selected
- ◉ Easy to implement and relatively low overhead
- ◉ Unfortunately, FIFO can replace heavily used pages which is the oldest

Oldest page is labeled in red

	2	3	2	1	5	2	4	5	3	2	5	2
	2	2	2	2	5	5	5	5	3	3	3	3
		3	3	3	3	2	2	2	2	2	5	5
				1	1	1	4	4	4	4	4	2
	F	F		F	F	F	F		F		F	F

- ◉ There are 9 faults with this reference string – the **hit rate is 25%**
- ◉ Not that practical for real-life systems

# Least-Recently-Used Replacement

- ◉ Exploits temporal locality by selecting the victim page that has **not been referenced for the longest time**
- ◉ Can provide better performance than FIFO
- ◉ Increased overhead
  - ◉ Hardware needs to maintain **the timings** of last reference of **all** pages
  - ◉ When replacing a page, OS **scans all time fields** to find the least-recently-used page
- ◉ LRU may perform poorly if the least-recently used page is the next page to be referenced by a program
  - ◉ e.g., a while loop may consist of many virtual pages, when jump back to the top, the **LRU page maybe the one that going to be referenced next**

# Least-Frequently-Used Replacement

- ◉ The victim page is the page that is the **least intensively referenced**
  - ◉ Based on the heuristic that a page not referenced often is not likely to be referenced in the future
- ◉ Each page has a **counter**, and is updated each time the page is referenced
- ◉ Have the possibility of selecting wrong page for replacement
  - ◉ A page that **was referenced heavily in the past** may never be referenced again, but will stay in memory while newer, **active pages are replaced**

# LRU & LFU

## ⦿ LRU

- ⦿ There are 7 faults with this reference string – the hit rate is 41.7%

Least-recently used pages are labeled in red

	2	3	2	1	5	2	4	5	3	2	5	2
	2	2	2	2	2	2	2	2	3	3	3	3
		3	3	3	5	5	5	5	5	5	5	5
				1	1	1	4	4	4	2	2	2
	F	F		F	F		F		F	F		

## ⦿ LFU

- ⦿ There are 6 faults with this reference string – the hit rate is 50%

Least-frequently used pages are labeled in red. Use FIFO to break the tie.

	2	3	2	1	5	2	4	5	3	2	5	2
	2	2	2	2	2	2	2	2	2	2	2	2
		3	3	3	5	5	5	5	5	5	5	5
				1	1	1	4	4	3	3	3	3
	F	F		F	F		F		F			

# Approximating LRU

- ◉ A way to **approximate LRU** with little overhead is to use a **use (reference) bit** to indicate that a page has recently been referenced
- ◉ The use bit of a virtual page is set to 1 by hardware
  - ◉ When the page is first loaded in memory upon page fault; or
  - ◉ When the page is referenced again
- ◉ OS needs some way to clear the use bit
  - ◉ One possible method
    - ◉ To screen out pages that are not actively used anymore, the system **periodically resets all the use bits to zero**
    - ◉ On the assumption that active pages will be referenced again in the near future

# Clock Replacement

- ◉ Another scheme which is similar to the approximating LRU
- ◉ System has **a pointer** (like a clock hand) points to the virtual page which is the “**oldest**” at this moment
- ◉ The system treats all the page frames as in a circular list
- ◉ When it is time to find a victim page, the system **checks the pointer**
  - ◉ if it points to a page **with use bit equals 0**, replace this one
  - ◉ Otherwise, **reset the use bit to 0** and **advance the pointer** to next virtual page
    - ◉ Reason: although the page is the “oldest”, it has been recently accessed; thus, give it a second chance and treat it as a “new” page
  - ◉ The process continues until a page with use bit equals to 0 is found

# Clock Replacement

- There are 8 faults with this reference string – the hit rate is 33.3%



Use bit is 1

2	3	2	1		5	2	4	5		3	2	5	2
2	2	2	2	2	5	5	5	5	5	3	3	3	3
	3	3	3	3	3	2	2	2	2	2	2	2	2
			1	1	1	1	4	4	4	4	4	5	5
F	F		F		F	F	F			F		F	

# Belady's Anomaly

- Does a larger memory cache always help?

Ref	A	B	C	D	A	B	E	A	B	C	D	E
1	A			D			E					+
2		B			A			+		C		
3			C			B			+		D	

Ref	A	B	C	D	A	B	E	A	B	C	D	E
1	A				+		E				D	
2		B				+		A				E
3			C						B			
4				D						C		

- “Stack-like property”
  - A cache of size  $k+1$  includes the contents of a cache of size  $k$ ; thus larger cache will be at least as good as smaller one.
- Belady's Anomaly
  - For policies without such a property, e.g., FIFO, Random, adding space to the cache may hurt the cache hit rate.



# Fetch Strategy

- ◉ Fetch Strategy

- ◉ When a process's address space is divided into pages and not all pages need to be in main memory. OS needs to consider when to load in a page?

- ◉ Demand Paging

- ◉ System **loads** a virtual page only when the running process **explicitly references** the page
  - ◉ Pro – Only loads pages that process actually need; space is not wasted
  - ◉ Con – Every time a new page is referenced, a page fault is generated, the process must wait

# Fetch Strategy

- ◉ Prefetching (Anticipatory Paging)
  - ◉ OS attempts to **predict** the virtual pages a process will need and **preloads** these pages **when has free page frames**
  - ◉ Must be carefully designed so that overhead incurred by the strategy does not reduce system performance
    - ◉ This strategy requires significant resources – page frames and disk I/O. If inaccurately determines which pages a process will need, might result in worse performance than in a demand paging system
- ◉ In Linux and other OSs, Demand + Prefetching to exploit spatial locality

# Thrashing

- ⦿ A **serious issue** appears in a multiprogramming or time sharing system using virtual memory technique
- ⦿ With **many processes** running, each **competes for physical memory** to place its virtual pages
- ⦿ If there are too many processes, **the memory demand exceeds the available physical memory**, what will be the consequence?
  - ⦿ Processes will be **busy swapping pages in and out**, and we see that the **page-fault rate** will be **very high**

# Thrashing

- ◉ Demands for page frames are too great, while one process is fetching a page and is waiting, the pages it already has can be stolen by other processes; when it resumes, it immediately faults again
- ◉ This is not just happening to one process; is **experienced by all**
- ◉ This leads to **low CPU utilization**
  - ◉ the processor is spending a significant amount of time doing nothing – all processes are waiting on page-in requests
- ◉ Solution: Suspend or kill some of the processes; don't allow users to oversubscribe

# Summary

- ◉ Virtual Memory system gives an illusion to the process that it has large amount of main memory to store the process's address space
- ◉ In real life, physical memory is scarce resource; OS makes use of slower, larger disks to support the virtualization of memory
- ◉ When the CPU tries to access a virtual page that is not in physical memory, OS will be invoked to handle this; it is responsible to load the page from swap space to main memory
- ◉ If the system does not have enough free physical memory, OS needs to make the decision in selecting some pages to swap out

# Summary

- ◉ Page replacement policy is critical to system performance; a wrong decision will induce more page faults
- ◉ Realistic replacement policies make use of past accessing history to guide the OS in selecting suitable pages for eviction
- ◉ LRU and LFU are performing better than others; however, they are more complicated and have higher implementation overhead
- ◉ Thrashing will appear if the system is oversubscribed with too many running processes

# Operating Systems

## ◉ Virtualization

- ◉ CPU Virtualization
  - ◉ Process Abstract
    - ◉ Address space
    - ◉ Process states
    - ◉ Process control block
    - ◉ Process operations API
    - ◉ Signals
  - ◉ Limited Direct Execution
    - ◉ System calls
    - ◉ Context switch
    - ◉ Interrupts
  - ◉ Scheduling
    - ◉ Scheduling metrics
    - ◉ FIFO, SJF, HRRN, STCF, RR, MLFQ
    - ◉ Multi-core scheduling, Linux CFS
- ◉ Memory Virtualization
  - ◉ Address space
  - ◉ Address translation: dynamic relocation
  - ◉ Segmentation
  - ◉ Paging
  - ◉ TLB
  - ◉ Multi-level paging
  - ◉ Inverted page table
  - ◉ Swap space
  - ◉ Page replacement policy: FIFO, LFR, LRU, Clock
  - ◉ Thrashing

## ◉ Concurrency

- ◉ Thread
  - ◉ POSIX threads (pthreads)
  - ◉ Race conditions, critical sections, mutual exclusion, atomic operations, synchronization
- ◉ Locks
  - ◉ Atomic instructions: test-and-set, compare-and-swap
  - ◉ Mutex locks
- ◉ Condition Variables
  - ◉ Pthread CVs
  - ◉ Producer-Consumer problem
- ◉ Semaphores
  - ◉ Binary Semaphores
  - ◉ Counting Semaphores
  - ◉ Ordering
  - ◉ Readers-Writers problem
- ◉ Deadlock
  - ◉ Dining philosophers' problem
  - ◉ Four necessary conditions
  - ◉ Deadlock prevention, avoidance, detection&recovery

## ◉ Persistence

- ◉ I/O devices (HDD, SSD)
- ◉ Files and Directories
  - ◉ Inode
  - ◉ File descriptor
  - ◉ Hard/Symbolic links
- ◉ File System Implementation
  - ◉ On-disk data structure
    - ◉ Superblock, Bitmap, Inodes, Data blocks
  - ◉ Free space management
    - ◉ Bitmap, linked-list, block-list
  - ◉ Caching and buffering
  - ◉ Access control and protection
  - ◉ Journaling file system
    - ◉ Data journaling
    - ◉ Metadata journaling
- ◉ **Advanced Topics**