

Files and Directories

2023 Fall COMP3230A

Contents

- ◉ Overview of storage disks
- ◉ Overview of file systems
- ◉ What is a file?
- ◉ What is a directory?

Related Learning Outcome

- ◉ ILO 2d - describe the principles and techniques used by OS to support persistent data storage

Readings & References


- ◉ Required Reading
 - ◉ Chapter 39 – **Interlude: Files and Directories**
 - ◉ <http://pages.cs.wisc.edu/~remzi/OSTEP/file-intro.pdf>
- ◉ References
 - ◉ Chapter 36 – **I/O Devices**
 - ◉ <http://pages.cs.wisc.edu/~remzi/OSTEP/file-devices.pdf>
 - ◉ Chapter 37 – **Hard Disk Drives**
 - ◉ <http://pages.cs.wisc.edu/~remzi/OSTEP/file-disks.pdf>
 - ◉ Chapter 44 – **Flash-based SSDs**
 - ◉ <http://pages.cs.wisc.edu/~remzi/OSTEP/file-ssd.pdf>
 - ◉ The SSD Anthology: Understanding SSDs and New Drives from OCZ
 - ◉ <http://www.anandtech.com/show/2738>
 - ◉ How do SSDs work?
 - ◉ <http://www.extremetech.com/extreme/210492-extremetech-explains-how-do-ssds-work>

Secondary Storage

- ◉ Most secondary storage devices involve magnetic disks, which are **random-access** storage
 - ◉ Data can be accessed by read-write head in any order
- ◉ Disk drives (magnetic or solid-state) are part of a class of storage called block devices.
 - ◉ These devices treat the storage space as a large 1-dimensional arrays of **logical disk blocks**, in which, the logical block is the **smallest unit of data transfer**
 - ◉ **Commonly-used disk block size is 4 KiB**
 - ◉ The logical blocks are addressed from 0 to $N-1$, where the disk has N logical disk blocks


Physical layout of HDD

- ◉ A disk consists of a number of magnetic platters with recording surfaces on both sides
- ◉ Rotate on spindle
- ◉ Each surface is divided into a number of concentric tracks
- ◉ Each track is divided into a number of sectors
- ◉ Vertical sets of tracks form cylinders
- ◉ Each surface has a disk head for reading and writing




Physical layout of HDD

- Each sector typically contains 512 bytes
- The unit of I/O operation is a **logical block**, typically of 4 KiB, which maps onto the sector(s)
 - Convert a logical block number into a cylinder #, a head #, and a sector #



HDD



Performance Characteristics of HDD


- ◉ The data in a particular disk sector can be read/written
- ◉ To access a data block
 - ◉ Disk arm must **move to the target track**; then **rotate** the disk to put **target sector** under the read-write head; then record is read-from/write-to the disk
- ◉ Performance characteristics
 - ◉ Seek time
 - ◉ Time for read-write head to move to target track from current location
 - ◉ average seek times is around 0.5 to 2 milliseconds
 - ◉ Rotational latency
 - ◉ Time for rotate the platter until the target sector is underneath read-write head
 - ◉ depends on the spinning rate; roughly around 2 ms
 - ◉ Transfer time
 - ◉ Time for further rotate the head to read/write the entire sector and transfer the data

Performance Considerations of HDD

- ◉ Ways to improve disk I/O performance
 - ◉ Disk scheduling
 - ◉ In multiprogramming environment, multiple processes can generate I/O requests at the same time, there may have several pending requests queued up at the disk queue
 - ◉ Which request should the system do first?
 - ◉ Because of the high cost of I/O, the OS historically played a role in deciding the order of I/Os issued to the disk
 - ◉ To optimize the data transfer with the minimum mechanical motion - seek time and rotational time
 - ◉ Caching
 - ◉ A disk cache buffer (in main memory) is used to temporarily hold disk data
 - ◉ Defragmentation
 - ◉ Place related data in contiguous sectors
 - ◉ Decreases number of seek operations required
 - ◉ Multiple disks
 - ◉ Disk I/O performance may be increased by spreading the operation over multiple disks

Physical layout of SSD


- SSD storage medium is called NAND **flash memory**
- Flash memory is non-volatile memory, which is organized as **a grid of storage cells**
 - Depends on the technologies, each cell can store 1 bit or 2 bits or 3 bits (or even more)
- A group of cells is organized into a “page”, which is **the smallest** structure that’s readable/ “writable” in a SSD
 - Today 4KiB (or 8KiB) pages are common on SSDs
- Pages are grouped together into blocks
 - It’s common to have **128 pages in a block** (**512KiB in a block**)



(Source: The SSD Anthology)

Physical layout of SSD

- ◉ Blocks are then grouped into planes, and you'll find multiple planes on a single NAND-flash chip
- ◉ A block is **the smallest** structure that can be erased in a NAND-flash device
 - ◉ You can read from and “write” to a page
 - ◉ But you cannot rewrite to a page unless **erasing the whole block** (128 pages at a time!!)
 - ◉ This is where many of the SSD's problems stem from
- ◉ Another issue is that frequently erase and write to a page/block will cause it to **wear out** (around 100,000 times)



(Source: The SSD Anthology)

Performance Characteristics of SSD

- ◉ Flash-based SSD provides the **standard block interface**
- ◉ To access a logical data block (e.g. 4 KiB)
 - ◉ The built-in control logic turns the requests into low-level **read, erase, and write** commands on the underlying physical blocks and physical pages
- ◉ Performance characteristics
 - ◉ Read a physical page
 - ◉ Able to access any location with the same performance - **random access device**
 - ◉ Typically quite fast, around **10s of microseconds**
 - ◉ Erase the whole block
 - ◉ It is quite expensive as it takes **a few milliseconds**
 - ◉ In addition, to preserve some data in the block, they must be copied to somewhere before the erase
 - ◉ Write (program) to a page
 - ◉ Usually takes around **100s of microseconds**

Block:	0				1				2			
Page:	00	01	02	03	04	05	06	07	08	09	10	11
Content:												

SSD Program/Erase: An Example

Page 0	Page 1	Page 2	Page 3
00011000	11001110	00000001	00111111
VALID	VALID	VALID	VALID

Page 0	Page 1	Page 2	Page 3
11111111	11111111	11111111	11111111
ERASED	ERASED	ERASED	ERASED

Page 0	Page 1	Page 2	Page 3
00000011	11111111	11111111	11111111
VALID	ERASED	ERASED	ERASED

HDD vs. SSD: Why do we still have HDD?

- ◉ Random I/Os: SSD >> HDD
- ◉ Sequential I/Os: SSD > HDD
- ◉ SSD: Random Reads < Random Writes?

Device	Random		Sequential	
	Reads (MB/s)	Writes (MB/s)	Reads (MB/s)	Writes (MB/s)
Samsung 840 Pro SSD	103	287	421	384
Seagate 600 SSD	84	252	424	374
Intel SSD 335 SSD	39	222	344	354
Seagate Savvio 15K.3 HDD	2	2	223	223

Abstraction of Persistent Storages

- ◉ Two key abstractions – Files and Directories
- ◉ What is a file?
 - ◉ From a **user's perspective**, it is a collection of related information that is **recorded on persistent storage** with a **human-readable name** given to it
 - ◉ From the **system's perspective**, it is a **linear array of bytes**, grouped in (logical) **blocks**, stored in somewhere, and has some kind of **low-level id** given to the file
 - ◉ In Unix systems, we call this low-level id – **inode number**
 - ◉ In Windows systems, it is called **file reference number**
 - ◉ This low-level id leads us to **a data structure**, where **the attributes of the files** are kept, e.g., locations of data of the file, ownership, etc.
- ◉ What is a directory?
 - ◉ Actually, **it is a file**, but its file content **is a mapping table** that maps filenames (in that directory) to their low-level ids
 - ◉ one entry for each file in that directory; can be a regular file or a directory file

File Systems

- ◉ Files (include directories) are managed by OS, and the part of OS dealing with files is known as the file system
 - ◉ File Management
 - ◉ Providing services to users and applications in the use of files & directories
 - ◉ Users should be able to refer to their files by symbolic names rather than having to use physical device names and physical location
 - ◉ Storage management
 - ◉ Allocating space for files on storage devices
 - ◉ File integrity
 - ◉ To guarantee, to the extent possible, that the data in the file are valid
 - ◉ Security
 - ◉ Data stored in file systems should be subject to strict access controls

File Abstraction

- ◉ From a user's standpoint, how to
 - ◉ locate the file, name the file, access the file, protect the file
- ◉ File system represents file as an **abstract data type**,
 - ◉ **which consists of a set of operations:**
 - ◉ Open – associate the target file to the process, allowing the process to perform specific functions on the file (returns a **file descriptor**)
 - ◉ Close – process no longer perform functions on the file until it is reopened
 - ◉ Create – a new file is defined (**space is allocated**) and **a new entry** must be made in the directory
 - ◉ Delete – release all file space and erase the directory entry
 - ◉ Write – make updates to the file according to the **current file pointer**
 - ◉ Read – copy data (starting from location points by the **current file pointer**) from a file to the memory
- ◉ Open file table
 - ◉ Each process maintains an array of file descriptors, each of which refers to an entry in the system-wide open file table.
 - ◉ Each entry tracks which underlying file the descriptor refers to, the current offset, and other relevant details such as whether the file is readable or writable.

File Abstraction

- ◉ ***which consists of a set of attributes (**metadata**) associate to a file:***

- ◉ Name – human-readable name
- ◉ Low-level id – unique tag **identifies** a file within file system
- ◉ Location – pointer to **storage locations of the data** of the file on device
- ◉ Size – current file size
- ◉ Accessibility – restrictions placed on access to file data
 - ◉ controls who can do Read, Write, Execute
- ◉ Time, date, and user identification – data for protection, security, and usage monitoring

⋮

- ◉ Where to store the metadata that associated to each file?

- ◉ Partly in the directory
- ◉ Mostly in the **file control block** (FCB) of that file
 - ◉ In Unix, this is the **inode**; in Windows, this is the **file record**


Directories

- ◉ As said, a directory is also a file
- ◉ We can view the content of a directory as a symbol table that **maps file/directory names** (in that directory) to the **directory entries** of that directory, which contain
 - ◉ the **low-level id** and a few other details of that file/directory
- ◉ Operations on directory
 - ◉ Search for a file
 - ◉ Create a file
 - ◉ Delete a file
 - ◉ List a directory
 - ◉ Rename a file
 - ◉ Traverse the file system

Name	Low-level ID
.	34
..	56
c0230a	123
c0234a	125

Directory Structure

- ◉ Hierarchically Structured File System
 - ◉ By placing directories within other directories, we have a directory tree, where all files and directories are stored
- ◉ A file system starts at a root directory “/”
 - ◉ The root directory contains various directories in the directory hierarchy
- ◉ The full name of a file is usually formed as the pathname from the root directory to the file – absolute path name
 - ◉ e.g., /home/c3230a/src/ws5.cc
 - ◉ Pros
 - ◉ File names need to be unique only within a given directory
 - ◉ Give more flexibility to users to name and group files
 - ◉ Efficient searching – by simply traverse the path to locate the files

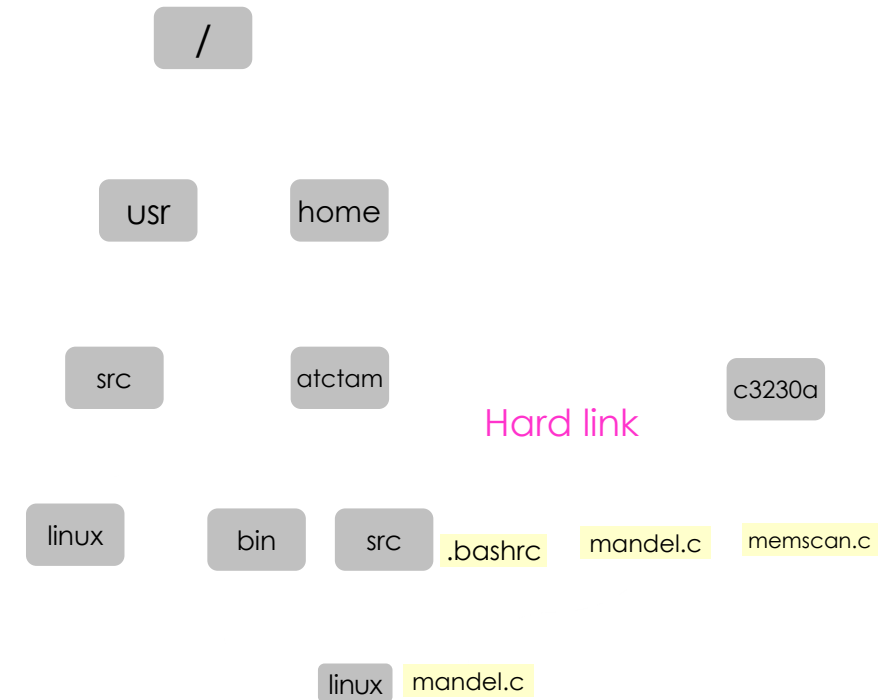


Directory Structure

- ◉ To simplify the navigation by using absolute path name, the concept of “Working directory” (**Current directory**) is used
 - ◉ Enables users to specify a pathname that does not begin at the root directory – **relative path name**
 - ◉ Absolute path (i.e., the path beginning at the root) = working directory + relative path name
- ◉ **Link**: a mechanism to **create another directory entry** that refers to an **existing** file/directory **in another location**
 - ◉ Adv: Facilitates data sharing and can make it easier for users to access files located throughout a file system’s directory structure

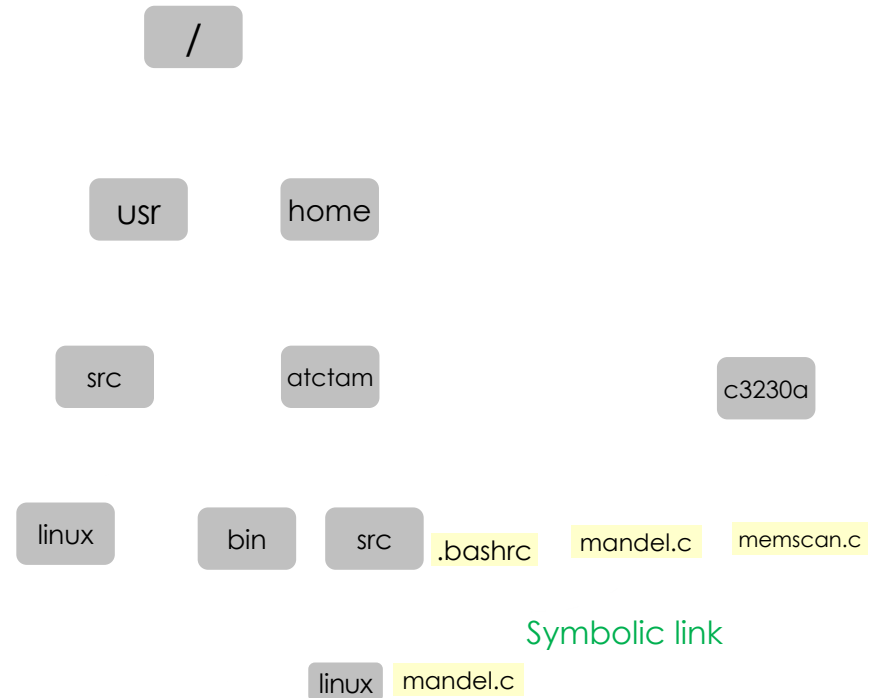
Directory Structure

- ◉ **Hard link**: create another directory entry that **maps to the same low-level id** of the original file
 - ◉ Unix – **ln** target new
 - ◉ Windows – CreateHardLink
- ◉ The file is not copied at all; the system just creates **two directory entries** at different locations **but refer to the same inode** (file control block)
- ◉ Remove one directory entry will not cause the file to be deleted
 - ◉ The system keeps track on how many different directory entries have been linked to the same low-level id – **reference count**
 - ◉ The system deletes the file, only when reference count reaches zero



Directory Structure

- ◉ Limitation of hard link
 - ◉ Can't create a hard link to directory
 - ◉ Can't create hard link to files in other disk partitions (i.e. another file system)
- ◉ Symbolic link: create another file that contains the pathname of original file as its data
 - ◉ Unix – **ln -s** target new
 - ◉ Windows – mklink, Shortcut
- ◉ Symbolic link is a special file type
- ◉ Remove the original file causes the soft link to be invalid – dangling reference



```
atctam@atctam-LinuxPC:~/src> ls -l linux
lrwxrwxrwx 1 atctam users 14 Nov 30 16:48 linux -> /usr/src/linux
atctam@atctam-LinuxPC:~/src> stat linux
  File: `linux' -> `/usr/src/linux'
  Size: 14      Blocks: 0      IO Block: 4096   symbolic link
```


Summary

- ◉ Describe what is the basic structure of a storage disk
- ◉ Discuss a few factors that affect the performance of the storage systems
- ◉ Understand the key concept of the file system – the FILE
 - ◉ What a file is consisted of? How a file provide persistent storage to the user?
- ◉ Describe what is the purpose of using DIRECTORY
 - ◉ What a directory is? How it is being structured? How it is related to the files?

Operating Systems

◉ Virtualization

- ◉ CPU Virtualization
 - ◉ Process Abstract
 - ◉ Address space
 - ◉ Process states
 - ◉ Process control block
 - ◉ Process operations API
 - ◉ Signals
 - ◉ Limited Direct Execution
 - ◉ System calls
 - ◉ Context switch
 - ◉ Interrupts
 - ◉ Scheduling
 - ◉ Scheduling metrics
 - ◉ FIFO, SJF, HRRN, STCF, RR, MLFQ
 - ◉ Multi-core scheduling, Linux CFS
- ◉ Memory Virtualization
 - ◉ Address space
 - ◉ Address translation: dynamic relocation
 - ◉ Segmentation
 - ◉ Paging
 - ◉ TLB
 - ◉ Multi-level paging
 - ◉ Inverted page table
 - ◉ Swap space
 - ◉ Page replacement policy: FIFO, LFR, LRU, Clock
 - ◉ Thrashing

◉ Concurrency

- ◉ Thread
 - ◉ POSIX threads (pthreads)
 - ◉ Race conditions, critical sections, mutual exclusion, atomic operations, synchronization
- ◉ Locks
 - ◉ Atomic instructions: test-and-set, compare-and-swap
 - ◉ Mutex locks
- ◉ Condition Variables
 - ◉ Pthread CVs
 - ◉ Producer-Consumer problem
- ◉ Semaphores
 - ◉ Binary Semaphores
 - ◉ Counting Semaphores
 - ◉ Ordering
 - ◉ Readers-Writers problem
- ◉ Deadlock
 - ◉ Dining philosophers' problem
 - ◉ Four necessary conditions
 - ◉ Deadlock prevention, avoidance, detection&recovery

◉ Persistence

- ◉ I/O devices (HDD, SSD)
- ◉ Files and Directories
 - ◉ Inode
 - ◉ File descriptor
 - ◉ Hard/Symbolic links
- ◉ File System Implementation
 - ◉ On-disk data structure
 - ◉ Superblock, Bitmap, Inodes, Data blocks
 - ◉ Free space management
 - ◉ Bitmap, linked-list, block-list
 - ◉ Caching and buffering
 - ◉ Access control and protection
 - ◉ Journaling file system
 - ◉ Data journaling
 - ◉ Metadata journaling
- ◉ *Advanced Topics*