

# Tutorial 3 - pthread

COMP3230 Principle of Operating System

Huang Songlin [huangs0@hku.hk](mailto:huangs0@hku.hk), 2023/10/12



# Table of Content

- POSIX Thread
- Race Condition & Solution
- Synchronization & Conditional Variable

# POSIX Thread

POSIX := Portable Operating System Interface of UNIX

- High-Level Design & Thread Lifecycle
- Creation
- Thread Function
- Termination & Collection

In code: **#include <pthread.h>**

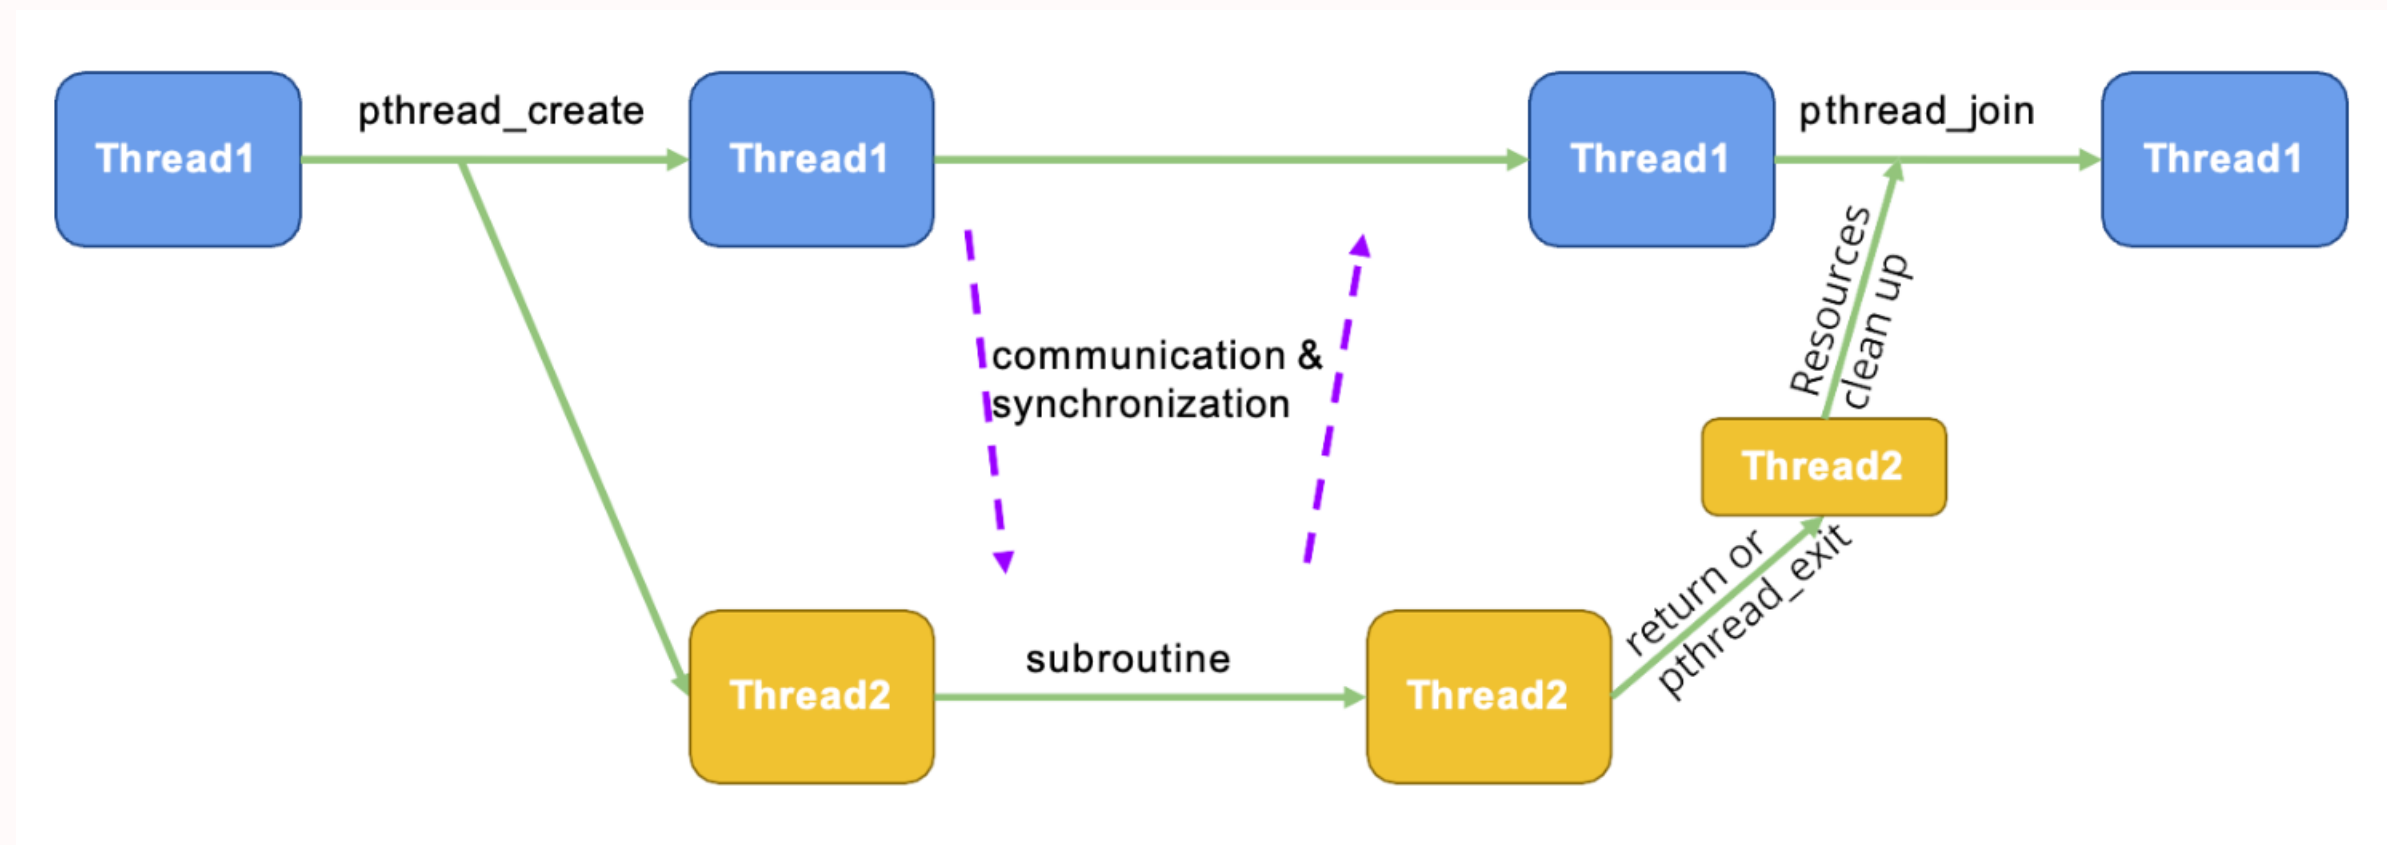
In compile: **-pthread** (gcc main.c -pthread)



# High-Level Design & Thread Lifecycle (Recap)

Threads is lightweight process, with:

- Shared heap (and global variables)
- Independent stack, register, and control (Thread Control Block)





# A Progressive Guide

Now you're given a task, multiply an array of 100 elements (0-99) with 2, a solution is:

```
1  #include <stdlib.h>
2  #include <stdio.h>
3
4  int main() {
5      int *a;
6      a = malloc(100 * sizeof(int));
7      for (int i = 0; i < 100; i++) { a[i] = i; }
8
9      for (int i = 0; i < 100; i++) {
10         a[i] = a[i] * 2;
11     }
12 }
```

- Computation is independent so multi-threading is a good option for acceleration.
- Multi-threading is like distributing tasks to workers (threads). For example, 2 threads:
  - Thread 1 → 0 ~ 50
  - Thread 2 → 50 ~ 100

# Define Thread Functions

Threads are like workers, multi-threading is like distributing tasks to workers. For example, Thread 1 → 0 - 50, Thread 2 → 50 - 100

```
1  #include <stdlib.h>
2  #include <stdio.h>
3
+ 4  int* a;
5
6  void* thr1_func(void* arg) {
7      for (int i = 0; i < 50; i++) { a[i] = a[i] * 2; }
8  }
9
10 void *thr2_func(void* arg) {
11     for (int i = 50; i < 100; i++) { a[i] = a[i] * 2; }
12 }
13
14 int main() {
- 15     int* a;
16     a = malloc(100 * sizeof(int));
17     for (int i = 0; i < 100; i++) { a[i] = i; }
18
19     for (int i = 0; i < 100; i++) {
20         a[i] = a[i] * 2;
21     }
22 }
```

- All thread function must have:
  - void pointer (**void \***) as input
  - void pointer (**void \***) as output

Questions:

- Both threads have **int i**, will they interrupt?
  - No, threads has **independent stack**.
- Why move line 17 to line 4?
  - If line 17, though **a** points to heap, **a** itself (the pointer) is on the stack of main function
  - If line 4, **a** become a global variable → all threads can access.

# Create Threads

Now it's time to create and run threads

```
1 #include <stdlib.h>
2 #include <stdio.h>
+ 3 #include <pthread.h>
4
5 int* a;
6
7 void* thr1_func(void* arg) {
8     for (int i = 0; i < 50; i++) { a[i] = a[i] * 2; }
9 }
10
11 void *thr2_func(void* arg) {
12     for (int i = 50; i < 100; i++) { a[i] = a[i] * 2; }
13 }
14
15 int main() {
16     a = malloc(100 * sizeof(int));
17     for (int i = 0; i < 100; i++) { a[i] = i; }
18
19     pthread_t thr1, thr2;
20     pthread_attr_t *attr = NULL;
21     int iret;
22
23     iret = pthread_create(&thr1, attr, thr1_func, NULL);
24     if (iret != 0) { perror("Cannot Create Thread 1"); }
25     iret = pthread_create(&thr2, attr, thr2_func, NULL);
26     if (iret != 0) { perror("Cannot Create Thread 2"); }
27
28     for (int i = 0; i < 100; i++) {
29         a[i] = a[i] * 2;
30     }
31 }
```

- Include header file **#include <pthread.h>**
- Define pthread\_t for each thread, respectively
  - **pthread\_t** is ID of thread.
- Define **pthread\_attr\_t** to configure thread behaviour
  - Mostly use NULL to use default

Call **pthread\_create**, 4 parameter:

- **pthread\_t \*thread\_id**
- **pthread\_attr\_t \*attr**
- **void \*(\*thr\_func)(void \*)**: Just a function whose input and output type both are void pointer. Functional Pointer (Not Required)
- **void \*arg**: the argument to be passed into **thr\_func**

Will return 0 if creation is success and non-zero instead

**Created threads start working immediately.**

# Terminate Threads

```
1 #include <stdlib.h>
2 #include <stdio.h>
3 #include <pthread.h>
4
5 int* a;
6
7 void* thr1_func(void* arg) {
8     for (int i = 0; i < 50; i++) { a[i] = a[i] * 2; }
9     pthread_exit(NULL);
10 }
11
12 void *thr2_func(void* arg) {
13     for (int i = 50; i < 100; i++) { a[i] = a[i] * 2; }
14     pthread_exit(NULL);
15 }
16
17 int main() {
18     a = malloc(100 * sizeof(int));
19     for (int i = 0; i < 100; i++) { a[i] = i; }
20
21     pthread_t thr1, thr2;
22     pthread_attr_t *attr = NULL;
23     int iret;
24
25     iret = pthread_create(&thr1, attr, thr1_func, NULL);
26     if (iret != 0) { perror("Cannot Create Thread 1"); }
27     iret = pthread_create(&thr2, attr, thr2_func, NULL);
28     if (iret != 0) { perror("Cannot Create Thread 2"); }
29
30     iret = pthread_join(thr1, NULL);
31     if (iret != 0) { perror("Cannot Collect Thread 1"); }
32     iret = pthread_join(thr2, NULL);
33     if (iret != 0) { perror("Cannot Collect Thread 2"); }
34 }
```

Similar to process, thread termination involves **two stages**:

1. Thread exit itself by **pthread\_exit**
2. Main thread collect value by **pthread\_join**

Thread function call **void pthread\_exit(void\* *retval*)** to exit:

- **void\* *retval***: return value to be collected by **pthread\_join**
- No return

Main thread call **int pthread\_join(pthread\_t *thr*, void\* *retval*)** to wait for termination:

- **pthread\_join will block main thread (like waitpid)**
- **pthread\_t *thr***: specify which thread to wait for
- **void\*\* *retval***: return value from **pthread\_exit**, use NULL if nothing to collect



# Improvement - Accepting Parameters

Defining functions for threads respectively is not extensible. A better way is to identify themselves using parameters

```
1  #include <stdlib.h>
2  #include <stdio.h>
3  #include <pthread.h>
4
5  int* a;
6
7  void* thr_func(void* arg) {
8      int id = * (int*) arg; // (int*) turn void ptr to int ptr
9      for (int i = 0; i < 50; i++) { a[i] = a[i] * 2; }
10     for (int i = 50 * id; i < 50 * (id + 1); i++) { a[i] = a[i] * 2; }
11     pthread_exit(NULL);
12 }
13
14 int main() {
15     a = malloc(100 * sizeof(int));
16     for (int i = 0; i < 100; i++) { a[i] = i; }
17
18     pthread_t thr1, thr2;
19     pthread_t * thrs = malloc(2 * sizeof(pthread_t));
20     pthread_attr_t *attr = NULL;
21     int iret;
22     int *args = malloc(2 * sizeof(int)); args[0] = 0; args[1] = 1;
23
24     iret = pthread_create(&thr1, attr, thr1_func, NULL);
25     if (iret != 0) { perror("Cannot Create Thread 1"); }
26     for (int i = 0; i < 2; i++) {
27         iret = pthread_create(&thrs[i], attr, thr_func, &args[i]);
28         if (iret != 0) { perror("Cannot Create Thread %d", i); }
29     }
30
31     for (int i = 0; i < 2; i++) {
32         iret = pthread_join(thrs[i], NULL);
33         if (iret != 0) { perror("Cannot Collect Thread 1"); }
34     }
35 }
```

We can use typecast on **void\* arg** to receive parameters:

```
int id = * (int*) arg;
```

And define pthread\_t(s) and argument(s) by **malloc**, and use the for loop to create threads.

Think question: Merge loop of creation and termination?

```
1  for (int i = 0; i < 2; i++) {
2      iret = pthread_create(&thrs[i], attr, thr_func, &args[i]);
3      if (iret != 0) { perror("Cannot Create Thread %d", i); }
4
5      iret = pthread_join(thrs[i], NULL);
6      if (iret != 0) { perror("Cannot Collect Thread 1"); }
7  }
```

Think question: More complicated parameters?

# Improvement - Collecting Return Value

Let each thread return its id

```
1 #include <stdlib.h>
2 #include <stdio.h>
3 #include <pthread.h>
4
5 int* a;
6
7 void* thr_func(void* arg) {
8     int id = *(int*) arg; // (int*) turn void ptr to int ptr
9     for (int i = 50 * id; i < 50 * (id + 1); i++) { a[i] = a[i] * 2; }
10    pthread_exit(NULL);
+11    pthread_exit((void *)&id); // similarly, return 0 if success
12 }
13
14 int main() {
15     a = malloc(100 * sizeof(int));
16     for (int i = 0; i < 100; i++) { a[i] = i; }
17
18     pthread_t * thrs = malloc(2 * sizeof(pthread_t));
19     pthread_attr_t *attr = NULL;
20     int iret, *ret_val;
21     int args[] = {0, 1};
22
23     for (int i = 0; i < 2; i++) {
24         iret = pthread_create(&thrs[i], attr, thr_func, &args[i]);
25         if (iret != 0) { perror("Cannot Create Thread"); }
26     }
27
28     for (int i = 0; i < 2; i++) {
29         iret = pthread_join(thrs[i], (void **) &ret_val);
30         if (iret != 0) { perror("Cannot Collect Thread"); }
31         printf("Thread %d returns %d\n", i, *ret_val);
32     }
33 }
```

To collect the return value:

- In thread function, use pointer to return value (like id here) as the parameter of **pthread\_exit**
- In main thread, initialize receiver, and use double pointer as the 2nd parameter of **pthread\_join**

When it runs:

```
1 >>> Thread 0 returns 1700461583
2 >>> Thread 1 returns -734169344
3
4 >>> Thread 0 returns -1317724401
5 >>> Thread 1 returns 1991890688
6
7 >>> segmentation fault
```

What happened?

# Collecting Return Value - Correct

Use Heap Variable - Won't be cleared after exit

```
1 #include <stdlib.h>
2 #include <stdio.h>
3 #include <pthread.h>
4
5 int* a;
6
7 void* thr_func(void* arg) {
8     int id = * (int*) arg; // (int*) turn void ptr to int ptr
9     for (int i = 50 * id; i < 50 * (id + 1); i++) { a[i] = a[i] * 2; }
10    pthread_exit((void *)&id); // similarly, return 0 if success
+11    int *ret = malloc(1 * sizeof(int)); // use heap variable
+12    *ret = id;
+13    pthread_exit((void *)ret); // similarly, return 0 if success
14 }
15
16 int main() {
17     a = malloc(100 * sizeof(int));
18     for (int i = 0; i < 100; i++) { a[i] = i; }
19
20     pthread_t * thrs = malloc(2 * sizeof(pthread_t));
21     pthread_attr_t *attr = NULL;
22     int iret, *ret_val;
23     int args[] = {0, 1};
24
25     for (int i = 0; i < 2; i++) {
26         iret = pthread_create(&thrs[i], attr, thr_func, &args[i]);
27         if (iret != 0) { perror("Cannot Create Thread"); }
28     }
29
30     for (int i = 0; i < 2; i++) {
31         iret = pthread_join(thrs[i], (void **) &ret_val);
32         if (iret != 0) { perror("Cannot Collect Thread"); }
33         printf("Thread %d returns %d\n", i, *ret_val);
34     }
35 }
```

Use pre-assigned global variable

```
1 #include <stdlib.h>
2 #include <stdio.h>
3 #include <pthread.h>
4
5 int* a, * ret_vals;
6
7 void* thr_func(void* arg) {
8     int id = * (int*) arg; // (int*) turn void ptr to int ptr
9     for (int i = 50 * id; i < 50 * (id + 1); i++) { a[i] = a[i] * 2; }
10    pthread_exit((void *)&id); // similarly, return 0 if success
+11    ret_vals[id] = id;
+12    pthread_exit(NULL);
13 }
14
15 int main() {
16     a = malloc(100 * sizeof(int));
17     for (int i = 0; i < 100; i++) { a[i] = i; }
18
19     pthread_t * thrs = malloc(2 * sizeof(pthread_t));
20     pthread_attr_t *attr = NULL;
21     int iret;
+22    ret_vals = malloc(2 * sizeof(int));
23    int args[] = {0, 1};
24
25    for (int i = 0; i < 2; i++) {
26        iret = pthread_create(&thrs[i], attr, thr_func, &args[i]);
27        if (iret != 0) { perror("Cannot Create Thread"); }
28    }
29
30    for (int i = 0; i < 2; i++) {
-31        iret = pthread_join(thrs[i], (void **) &ret_val);
+32        iret = pthread_join(thrs[i], NULL);
33        if (iret != 0) { perror("Cannot Collect Thread"); }
-34        printf("Thread %d returns %d\n", i, *ret_val);
+35        printf("Thread %d returns %d\n", i, ret_vals[i]);
36    }
37 }
```

Recap question: Why we can use stack variable of main function to pass in parameter?

# Identification

For OS to manage threads, 1st is to identify them

```
1 #include <stdlib.h>
2 #include <stdio.h>
3 #include <pthread.h>
4 #include <sys/syscall.h>
5 #include <unistd.h>
6
7 void* thr_func(void* arg) {
8     int i = * (int*) arg; // (int*) turn void ptr to int ptr
9     printf("Thread %d: \nMy pthread id is %d \nMy process id is %d\n\n",
10           i, pthread_self(), getpid(), syscall(SYS_gettid));
11     pthread_exit(NULL);
12 }
13
14 int main() {
15
16     pthread_t * thrs = malloc(2 * sizeof(pthread_t));
17     pthread_attr_t *attr = NULL;
18     int iret, args[] = {0, 1};
19
20     printf("Main Thread Process id %d", getpid());
21
22     for (int i = 0; i < 2; i++) {
23         iret = pthread_create(&thrs[i], attr, thr_func, &args[i]);
24         printf("Main Thread: \nCreate Thread %d with pthread_id %d\n\n", i, thrs[i]);
25     }
26
27     for (int i = 0; i < 2; i++) {
28         iret = pthread_join(thrs[i], NULL);
29         if (iret != 0) { perror("Cannot Collect Thread"); }
30     }
31 }
```

Normally threads has:

- Thread ID: obtain by **pthread\_self()**, used by `<pthread.h>`, the same as `pthread_create(pthread* thread_id)`
- System ID: obtain by **syscall(SYS\_gettid)**, used by OS, the same form as PID in Linux (Linux thread ~= process)
- Process ID: obtain by **getpid()**, all threads within process share the same pid
- (Manual) ID: pass in by the arguments for customization.

```
1 >>> Main Thread Process id 3078547
2 >>> Main Thread:
3 >>> Create Thread 0 with pthread_id 527197952
4 >>>
5 >>> Main Thread:
6 >>> Create Thread 1 with pthread_id 518805248
7 >>>
8 >>> Thread 0:
9 >>> My pthread id is 527197952
10 >>> My process id is 3078547
11 >>> My system id is 3078548
12 >>>
13 >>> Thread 1:
14 >>> My pthread id is 518805248
15 >>> My process id is 3078547
16 >>> My system id is 3078549
```



# Cancellation

How to terminate (incorperative) thread from other thread?

```
1 #include <stdio.h>
2 #include <unistd.h>
3 #include <pthread.h>
4
5 void *thr1_func (void *arg){
6     printf("func1 disable cancellation\n");
7     pthread_setcancelstate(PTHREAD_CANCEL_DISABLE, NULL);
8     sleep(5);
9     printf("func1 enable cancellation\n");
10    pthread_setcancelstate(PTHREAD_CANCEL_ENABLE, NULL);
11    pthread_setcanceltype(PTHREAD_CANCEL_ASYNCHRONOUS, NULL);
12    while (1) { printf("..."); }
13    // NO EXIT
14 }
15
16 void *thr2_func (void *arg){
17     pthread_t tid = * (pthread_t *)arg;
18     for (int i = 0; i < 4; i++) {
19         printf("func2 sends a cancellation request to func1\n");
20         pthread_cancel(tid);
21         sleep(2);
22     }
23     pthread_exit(NULL);
24 }
25
26 int main() {
27     pthread_t thread_id1, thread_id2;
28
29     pthread_create(&thread_id1, NULL, thr1_func, NULL);
30     sleep(2);
31     pthread_create(&thread_id2, NULL, thr2_func, (void*)&thread_id1);
32
33     pthread_join(thread_id1, NULL);
34     printf("func1 has terminated\n");
35     pthread_join(thread_id2, NULL);
36     printf("func2 has terminated\n");
37     return 0;
38 }
```

`pthread_cancel(pthread_t thread_id)` can cancel threads from other threads like Thread 2 or main. (Like kill API for process)

Threads can configure their behavior on cancellation by:

- `pthread_setcancelstate(int state, int* old_state)`, states:
  - **PTHREAD\_CANCEL\_DISABLE**: ignore cancellation
  - **PTHREAD\_CANCEL\_ENABLE**: enable cancellation
  - `int* old_state`: for recording old state, mostly NULL
- `pthread_setcanceltype(int type, int* old_type)`, type:
  - **PTHREAD\_CANCEL\_DEFERRED**:  
cancellation will be deferred until the thread next call a function that is a cancellation point. (Default)
  - **PTHREAD\_CANCEL\_ASYNCHRONOUS**:  
cancel at any time and mostly immediately. (for manually cancel)

Threads cancelled is terminated, and shall also be collected.

# Takeaway - Comparison with process

In high-level, process models an instance of program, while threads are just stream of instructions, but they're highly similar.

|                   | Process  | Thread  |
|-------------------|--|---|
| Creation          | <code>fork()</code> → create heap + stack + PCB              | <code>pthread_create()</code> → create stack + TCB                    |
| Termination       | <code>exit()</code> → clear all resources                    | <code>pthread_exit()</code> → only clear stack, heap alive            |
| Force Termination | <code>kill()</code> / signal handler                         | <code>pthread_cancel()</code> / <code>pthread_setcancelstate()</code> |
| Collection        | <code>waitid()</code> / <code>waitpid()</code> — collect PCB | <code>pthread_join()</code> — collect value                           |
| Communication     | IPC — Signal/Pipe/Socket                                     | Shared Memory — heap / global variable                                |

# Race Condition & Solution

- Race Condition — Multiple threads accessing same memory
- Atom Operation — Guarantee provided by hardware
- Mutex Lock
  - Create & Destroy
  - Lock & Unlock

# Race Condition

Race condition arises if **multiple threads simultaneously read/write the shared resources (memory)** → **Unpredicted Outcome**

```
1 #include <stdlib.h>
2 #include <stdio.h>
3 #include <pthread.h>
4
5 int a = 0, iret;
6
7 void* thr_func(void* arg) {
8     for (int i = 0; i < 1000000; i++) { a += 1; }
9     pthread_exit(NULL);
10 }
11
12 int main() {
13     pthread_t thr1, thr2;
14
15     iret = pthread_create(&thr1, NULL, thr_func, NULL);
16     if (iret != 0) { perror("Cannot Create Thread"); }
17     iret = pthread_create(&thr2, NULL, thr_func, NULL);
18     if (iret != 0) { perror("Cannot Create Thread"); }
19
20     iret = pthread_join(thr1, NULL);
21     if (iret != 0) { perror("Cannot Collect Thread"); }
22     iret = pthread_join(thr2, NULL);
23     if (iret != 0) { perror("Cannot Collect Thread"); }
24
25     printf("%d", a);
26 }
```

Correct: 2000000, Actual: 1016738, 1003869, 1019114, ...

Question: Is the previous example race condition?

Reason: most C operations like `a += 1` involves:

```
1 thr_func:
2 .LFB6:
3     .cfi_startproc
4     ...
5     jmp .L2
6 .L3:
7     movl    a(%rip), %eax
8     addl    $1, %eax
9     movl    %eax, a(%rip)
10    addl    $1, -4(%rbp)
```

|                         |                                 |
|-------------------------|---------------------------------|
| Read Memory to Register | <code>movl a(%rip), %eax</code> |
|-------------------------|---------------------------------|

|                          |                             |
|--------------------------|-----------------------------|
| Do Operation on Register | <code>addl \$1, %eax</code> |
|--------------------------|-----------------------------|

|                          |                                 |
|--------------------------|---------------------------------|
| Write Register to Memory | <code>movl %eax, a(%rip)</code> |
|--------------------------|---------------------------------|

Thus, if Thread 1 and Thread 2 **read memory simultaneously**:

- They add 1 to their own register
- Then write back `a + 1` to `a`
- `a` in memory will be `a + 1` instead of `a + 2`



# Atomic Operation

As discussed in lecture, the correct way is to allow at most one thread access each shared data

For this naive case, we can use the **fetch\_and\_add** atomic:

```
1 int FetchAndAdd(int* ptr) {
2     int old = *ptr;
3     *ptr = old + 1;
4     return old;
5 }
```

🤔 Looks no diff w.r.t. previous code, still read/write memory.

🧐 YES, above is just illustration, following is actual code:

```
1 #include <stdlib.h>
2 #include <stdio.h>
3 #include <pthread.h>
4
5 int a = 0, iret;
6
7 void* thr_func(void* arg) {
8     for (int i = 0; i < 1000000; i++) {
9         a += 1;
10        __atomic_fetch_add(&a, 1, __ATOMIC_SEQ_CST);
11    }
12    pthread_exit(NULL);
13 }
14
15 int main() { ... }
```

Now the code will always leads to 2000000, correct answer.

What's the difference? Compare the assembly code:

With atomic\_fetch\_add

```
1 thr_func:
2 .LFB6:
3 .cfi_startproc
4 ...
5 jmp .L2
6 .L3:
7 lock addl $1, a(%rip)
8 addl $1, -4(%rbp)
```

Without  
atomic\_fetch\_add

```
1 thr_func:
2 .LFB6:
3 .cfi_startproc
4 ...
5 jmp .L2
6 .L3:
7 movl a(%rip), %eax
8 addl $1, %eax
9 movl %eax, a(%rip)
10 addl $1, -4(%rbp)
```

Correctness of Atomic is **guaranteed by hardware (CPU)** and requires the hardware support, not related with OS.

Other atomic operations supported in GCC(Linux):

- `__atomic_add_fetch`
- `__atomic_test_and_set`
- `__atomic_compare_exchange`

Reference: [GCC Atomic Built-in](#)

# Mutex := Mutually Exclusive

Atomic only cover simple cases, for more complicated scenario we need to **lock the critical section**

```
1 #include <stdlib.h>
2 #include <stdio.h>
3 #include <pthread.h>
4
5 int a = 0, iret;
+6 pthread_mutex_t lock;
7
8 void* thr_func(void* arg) {
9     for (int i = 0; i < 1000000; i++) {
+10 pthread_mutex_lock(&lock);
11         a++;
+12 pthread_mutex_unlock(&lock);
13     }
14     pthread_exit(NULL);
15 }
16
17 int main() {
18     pthread_t thr1, thr2;
+19 pthread_mutex_init(&lock, NULL); // NULL -> default
20
21     iret = pthread_create(&thr1, NULL, thr_func, NULL);
22     if (iret != 0) { perror("Cannot Create Thread"); }
23     iret = pthread_create(&thr2, NULL, thr_func, NULL);
24     if (iret != 0) { perror("Cannot Create Thread"); }
25
26     iret = pthread_join(thr1, NULL);
27     if (iret != 0) { perror("Cannot Collect Thread"); }
28     iret = pthread_join(thr2, NULL);
29     if (iret != 0) { perror("Cannot Collect Thread"); }
30
+31 pthread_mutex_destroy(&lock);
32
33     printf("%d", a);
34 }
```

- Before init, define the lock variable `pthread_mutex_t`.
- Initialize lock by `pthread_mutex_init` with two parameter:
  - `pthread_mutex_t mutex`: lock variable to be initialized
  - `pthread_mutexattr_t attr`: configuration, usually NULL
- Acquire the lock by `pthread_mutex_lock`
  - **IMPORTANT: `pthread_mutex_lock` is a blocked call, thread will be blocked until lock is acquired**
  - `pthread_mutex_t mutex`: lock to acquire
- Release the lock by `pthread_mutex_unlock`
  - `pthread_mutex_t mutex`: lock to acquire
- Finally, destroy the lock by `pthread_mutex_destroy`
  - `pthread_mutex_t mutex`: lock to acquire

Now the code will always leads to 2000000, correct answer.

Mutex lock gives more flexibility because you can do more between `pthread_mutex_lock` and `pthread_mutex_unlock`

# trylock

pthread\_mutex\_lock block threads until lock acquired, but threads can work on other tasks with no race condition.

Want thread to work on another task b if a is occupied.

```
1  #include <stdlib.h>
2  #include <stdio.h>
3  #include <pthread.h>
4
5  int a = 0, b = 0, iret;
6  pthread_mutex_t lock_a, lock_b;
7
8  void* thr_func(void* arg) {
9      for (int i = 0; i < 1000000; i++) {
10         if (pthread_mutex_trylock(&lock_a) == 0) {
11             a++;
12             pthread_mutex_unlock(&lock_a);
13         } else if (pthread_mutex_trylock(&lock_b) == 0) {
14             b++;
15             pthread_mutex_unlock(&lock_b);
16         }
17     }
18     pthread_exit(NULL);
19 }
20
21 int main() { ... }
22 >>> a: 1029395, b: 970605
```

Sum of a and b is 2000000

**pthread\_mutex\_lock** will not be appropriate as it'll block thread until lock acquired.

Instead, **pthread\_mutex\_trylock(pthread\_mutex\_t\* mutex)**:

- **Immediately** return 0 if lock acquired and non-zero if not
- **pthread\_mutex\_t mutex**: lock try to acquire

**pthread\_mutex\_trylock will not block thread**, if lock is not acquired, thread can do other things → more efficient usage.

# Synchronization & Condition Variable

- Synchronization — Controlling Execution Order
- Conditional Variable





# Synchronization — Controlling Execution Order

Threads are executed independently (and in parallel), but we do need the control on execution order.

For example, in example of  $0-99 \times 2$ , need to know if Thread 1 and Thread 2 have finished their task → **pthread\_join**

Question: a pthread\_join but thread keep alive? i.e., just wait.

```
1 #include <stdlib.h>
2 #include <stdio.h>
3 #include <pthread.h>
4 #include <unistd.h>
5
6 int* a, *state;
7 pthread_mutex_t lock; // not race, just for cond
8 pthread_cond_t cond;
9
10 void* thr_func(void* arg) {
11     int id = *(int*) arg; // (int*) turn void ptr to int ptr
12     for (int i = 50 * id; i < 50 * (id + 1); i++) { a[i] = a[i] * 2; }
13     pthread_mutex_lock(&lock);
14     state[id] = 1;
15     pthread_mutex_unlock(&lock);
16     pthread_cond_signal(&cond);
17     // can do more here
18     sleep(3);
19     pthread_exit(NULL);
20 }
21
22 int main() {
23     a = malloc(100 * sizeof(int));
24     for (int i = 0; i < 100; i++) { a[i] = i; }
25
26     pthread_mutex_init(&lock, NULL); // NULL -> default
27     pthread_cond_init(&cond, NULL); // NULL -> default
28
29     pthread_t * thrs = malloc(2 * sizeof(pthread_t));
30     pthread_attr_t *attr = NULL;
31     int iret, args[] = {0, 1};
32     state = malloc(2 * sizeof(int)); state[0] = 0; state[1] = 0;
33
34     for (int i = 0; i < 2; i++) {
35         iret = pthread_create(&thrs[i], attr, thr_func, &args[i]);
36     }
37
38     pthread_mutex_lock(&lock);
39     while (state[0] && state[1]) {
40         pthread_cond_wait(&cond, &lock);
41     }
42
43     printf("All threads have done their job and still alive"); fflush(stdout);
44     sleep(3);
45
46     for (int i = 0; i < 2; i++) {
47         iret = pthread_join(thrs[i], NULL);
48     }
49
50     pthread_mutex_destroy(&lock);
51     pthread_cond_destroy(&cond);
52 }
```

## Recap - Cond Variable

Cond variable involves an inner "queue", with two function:

- cond\_wait: wait for cond, **put myself in "queue" and blocked**
- cond\_signal: **wake up one thread in the "queue"** to work

## Usage

**IMPORTANT: Conditional variable must be used with mutex\_lock (especially cond\_wait) to protect thread safety**

Similar to mutex\_lock, we need to initialize **pthread\_cond\_t**:

1. Define the cond variable **pthread\_cond\_t cond**;
2. Do initialization by **pthread\_cond\_init(&cond, NULL)**
3. Finally destroy by **pthread\_cond\_destroy(&cond)**

Using cond variable involves two end:

- **pthread\_cond\_signal(&cond)**: wake up one thread in queue, no mutex protect need
- **pthread\_cond\_wait(&cond, &lock)**: (blocked) wait for a specific condition:
  - **Can only be called with mutex lock (&lock) acquired**
  - **Will release the mutex lock (&lock) once called**
  - **Will get back mutex lock once being waked up** (implicitly do pthread\_mutex\_lock)

# Cond Variable - What's state here?

Other than pthread\_cond\_t, there's an `int * state`, but why?

```
1 int* a, *state;
2 pthread_mutex_t lock; // not race, just for cond
3 pthread_cond_t cond;
4
5 void* thr_func(void* arg) {
6     int id = *(int*) arg; // (int*) turn void ptr to int ptr
7     for (int i = 50 * id; i < 50 * (id + 1); i++) { a[i] = a[i] * 2; }
8     pthread_mutex_lock(&lock);
9     state[id] = 1;
10    pthread_mutex_unlock(&lock);
11    pthread_cond_signal(&cond);
12    // can do more here
13    sleep(3);
14    pthread_exit(NULL);
15 }
16
17 int main() {
18     a = malloc(100 * sizeof(int));
19     for (int i = 0; i < 100; i++) { a[i] = i; }
20
21     pthread_mutex_init(&lock, NULL); // NULL -> default
22     pthread_cond_init(&cond, NULL); // NULL -> default
23
24     pthread_t * thrs = malloc(2 * sizeof(pthread_t));
25     pthread_attr_t *attr = NULL;
26     int iret, args[] = {0, 1};
27     state = malloc(2 * sizeof(int)); state[0] = 0; state[1] = 0;
28
29     for (int i = 0; i < 2; i++) {
30         iret = pthread_create(&thrs[i], attr, thr_func, &args[i]);
31     }
32
33     pthread_mutex_lock(&lock);
34     while (state[0] && state[1]) {
35         pthread_cond_wait(&cond, &lock);
36     }
37
38     printf("All threads have done their job and still alive"); fflush(stdout);
39     sleep(3);
40
41     for (int i = 0; i < 2; i++) {
42         iret = pthread_join(thrs[i], NULL);
43     }
44
45     pthread_mutex_destroy(&lock);
46     pthread_cond_destroy(&cond);
47 }
```

- Initialize the **state** to be `[0, 0]`
- Each thread **update their state before signal** main thread
  - States are shared among threads → need mutex\_lock
- Waited thread (main) **will check state once they wake up**
  - Recall cond\_wait acquire mutex\_lock when wake up
  - If state matches, go to next stage

Reason: **Cond Variable is "Stateless"!**

- Only identify an even happened (like a thread finish task)
- In simple cases like sync with only one thread,
- solely use cond variable is fine → one signal, one wait
- But more complicated cases, need help of shared variables like **state** → thus need mutex\_lock

Question: Replace while by if?

```
1 pthread_mutex_lock(&lock);
2 while (state[0] && state[1]) {
3     pthread_cond_wait(&cond, &lock);
4 }
```