

# Tutorial 4 - Semaphore

COMP3230 Principle of Operating System

Huang Songlin (huangs0@hku.hk), 2023-10-26

An abstract background on the left side of the slide, featuring flowing, wavy lines in various shades of orange and red, creating a sense of movement and depth.

# Agenda

1. Recap Tutorial 3
2. Semaphore vs Conditional Variable
3. Multi-Threading Debugging

# Recap Tutorial 3

- POSIX Thread (pthread.h)
  - thr\_func, create, join, exit
- Atomic & Mutex Lock
- Conditional Variable to be covered later

# POSIX Thread

Reminder: In code: `#include <pthread.h>` / In compile: `gcc -pthread` (or functions not found)

Thread Function: `void *thr_func(void *arg)`

Question: Difference Between `void` and `void*`?

Thread Creation: `pthread_create(pthread_t* id, NULL, void *thr_func(void *arg), void* arg)`

Question: What's `pthread_t`?

Thread Exit: `pthread_exit(void* exit_val)`

Question: What if `int a = 1` and `pthread_exit((void*) a);`?

Thread Termination: `pthread_join(pthread_t id, (void**) exit_val)`: **BLOCKED CALL**

Practical Question:

- Which scope can thread access?

Answer:

- Which scope independent for each thread?

Answer:

- How thread collects parameter and return?

Answer:

# Atomic & Mutex Lock

Race Condition: **multiple threads simultaneously read/write the shared resources (memory)**

```
5  int a = 0, iret;  
6  
7  void* thr_func(void* arg) {  
8      for (int i = 0; i < 1000000; i++) { a += 1; }  
9      pthread_exit(NULL);  
10 }  
11  
12 int main() {  
13     pthread_t thr1, thr2;  
14  
15     iret = pthread_create(&thr1, NULL, thr_func, NULL);  
16     if (iret != 0) { perror("Cannot Create Thread"); }  
17     iret = pthread_create(&thr2, NULL, thr_func, NULL);  
18     if (iret != 0) { perror("Cannot Create Thread"); }
```

Correct: 2000000, Actual: 1016738

Soln: one and only one can use shared resource

Basic: **Atomics** like `__atomic_add_fetch`

Hardware support (Lock Memory / Cache)

Mutex Locks - More customisable soln

1. Creation:

1. Create variable by `pthread_mutex_t lock`;
2. Init by `pthread_mutex_init(&lock, NULL)`;

2. Lock & Unlock


1. **BLOCKED** Lock by `pthread_mutex_lock(&lock)`;
2. **UNBLOCKED** Lock (return 0 if lock) by `pthread_mutex_trylock(&lock)`;
3. Unlock by `pthread_mutex_unlock(&lock)`;

3. Destroy by `pthread_mutex_destroy(&lock)`;



# Semaphore

- `<semaphore.h>` API
  - Unnamed / Named - Init / Destroy
  - Wait / Post
- Case Study: Semaphore vs Cond Variable



# Semaphore API - Unnamed

Reminder: in code `#include <semaphore.h>` / in compile `gcc -pthread`

```
1 #include <stdlib.h>
2 #include <stdio.h>
3 #include <pthread.h>
+ 4 #include <semaphore.h>
5 #include <unistd.h>
6
7 int* a;
+ 8 sem_t sem;
9
10 void* thr_func(void* arg) {
11     for (int i = 0; i < 100; i++) { a[i] = a[i] * 2; }
12     sem_post(&sem);
13     sleep(3); // Do customization here
14     pthread_exit(NULL);
15 }
16
17 int main() {
18     a = malloc(100 * sizeof(int));
19     for (int i = 0; i < 100; i++) { a[i] = i; }
20
21     sem_init(&sem, 0, 0);
22     pthread_t thr;
23     pthread_create(&thr, NULL, thr_func, NULL);
24
25     sem_wait(&sem);
26     printf("All threads finished job and still alive\n");
27
28     pthread_join(thr, NULL);
29
30     sem_destroy(&sem);
31 }
```

## Create Semaphore

First create variable `sem_t sem`; Then init by:

`int sem_init(sem_t* sem, int pshared, int value)`

- `sem_t* sem`: pointer of semaphore to be init
- `int pshared`: indicates if the semaphore is:
  - Local to this process (`pshared=0`)
  - shared among processes (`pshared=1`)
- `unsigned int value`: the initial value of sem

## Destroy Semaphore

`int sem_destroy(sem_t* sem)`

- `sem_t* sem`: pointer of sem to be destroyed

P.S. Return 0 if succeed and non-zero if failed

# Semaphore API - Post / Wait

Two Atomic Operation: Wait: P(&sem), sem\_wait(&sem) / Post: V(&sem), sem\_post(&sem)

```
1 #include <stdlib.h>
2 #include <stdio.h>
3 #include <pthread.h>
4 #include <semaphore.h>
5 #include <unistd.h>
6
7 int* a;
8 sem_t sem;
9
10 void* thr_func(void* arg) {
11     for (int i = 0; i < 100; i++) { a[i] = a[i] * 2; }
12     sem_post(&sem);
13     sleep(3); // Do customization here
14     pthread_exit(NULL);
15 }
16
17 int main() {
18     a = malloc(100 * sizeof(int));
19     for (int i = 0; i < 100; i++) { a[i] = i; }
20
21     sem_init(&sem, 0, 0);
22     pthread_t thr;
23     pthread_create(&thr, NULL, thr_func, NULL);
24
25     sem_wait(&sem);
26     printf("All threads finished job and still alive\n");
27
28     pthread_join(thr, NULL);
29
30     sem_destroy(&sem);
31 }
```

Wait: P(&sem)

`int sem_wait(sem_t* sem)`, try `sem -= 1`, and:

- if `sem > 0`, perform `sem -= 1` and return
- if `sem = 0`, blocked until `sem -= 1` can be performed (`sem > 0`) then do `sem -= 1` return

Post: V(&sem)

`int sem_post(sem_t* sem)`, do `sem += 1`, and:

- if no thread is waiting, simply return
- if some thread(s) is waiting, wake up one(to do `sem-=1` and return,so exactly one thread)



# Comparison with Cond Variable

Seems `pthread_cond_wait`  $\approx$  `sem_wait`, and `pthread_cond_signal`  $\approx$  `sem_post`? Hint: `sem = 0` is "signal"

`int pthread_cond_wait(pthread_cond_t*cond, ...)`

Being blocked until receiving signal of *cond*

`int pthread_cond_signal(pthread_cond_t*cond)`

Wake up one thread in waiting queue and  
do nothing if queue empty

`int sem_wait(sem_t* sem)`

if *sem* <= 0 Being blocked ,  
if *sem* > 0 do *sem* -= 1 and return

`int sem_post(sem_t* sem)`

Do *sem* += 1 and Wake up one waiting thread  
and do nothing if queue empty (*sem* > 0  
before)

Generally, Semaphore is Conditional Variable + an integer state.

# Case Study - no (int) state means what?

Let the only (child) thread do sth and inform parent work is done but not pthread\_join?

Semaphore Implementation

```
1  int* a;
2  sem_t sem;
3
4  void* thr_func(void* arg) {
5      for (int i = 0; i < 100; i++) { a[i] = a[i] * 2; }
+6  sem_post(&sem); // sem_post to say work done
7      pthread_exit(NULL);
8  }
9
10 int main() {
11     a = malloc(100 * sizeof(int));
12     for (int i = 0; i < 100; i++) { a[i] = i; }
13
14     sem_init(&sem, 0, 0);
15     pthread_t thr;
16     pthread_create(&thr, NULL, thr_func, NULL);
17
+18 sem_wait(&sem); // wait until job finished
19
20     pthread_join(thr, NULL);
21     sem_destroy(&sem);
22 }
```

Replacing sem\_post and sem\_wait:

```
1  int* a;
2  pthread_cond_t cond;
3
4  void* thr_func(void* arg) {
5      for (int i = 0; i < 100; i++) { a[i] = a[i] * 2; }
+6  pthread_cond_signal(&cond); // inform job finished
7      pthread_exit(NULL);
8  }
9
10 int main() {
11     a = malloc(100 * sizeof(int));
12     for (int i = 0; i < 100; i++) { a[i] = i; }
13
14     pthread_cond_init(&cond, NULL);
15     pthread_t thr;
16     pthread_create(&thr, NULL, thr_func, NULL);
17
+18 pthread_cond_wait(&cond); // wait until finished
19
20     pthread_join(thr, NULL);
21     pthread_cond_destroy(&cond);
22 }
```

# Case Study (Contd) Different Order?

Main thread and child thread executed simultaneously → order between post and wait is arbitrary.

If wait before post?

1. Main Thread call `sem_wait(&sem)`, try `sem -= 1` → find `sem = 0`, being blocked
2. Child Thread call `sem_post(&sem)`, do `sem += 1` → `sem = 1` → wake up main
3. Main Thread being waked up by child, finish `sem -= 1` → `sem = 0` → all done 🥳

1. Main Thread call `cond_wait(&cond)`, put itself in waiting queue and being blocked
2. Child Thread call `sem_post(&cond)`, wake up one from queue → wake up main
3. Main Thread being waked up and return → all done 🥳

If post before wait?

1. Child Thread call `sem_post(&sem)`, do `sem += 1` → `sem = 1` → queue is empty, do nothing and return
2. Main Thread call `sem_wait(&sem)`, try `sem -= 1` → now `sem = 1 > 0` → do `sem -= 1` and return → all done 🥳

1. Child Thread call `cond_signal(&cond)`, try to wake up one from waiting queue → queue is empty, do nothing and return
2. Main Thread call `cond_wait(&cond)`, put itself in waiting queue and being blocked → **No body will wake it up** 😞

Problem of Standalone Conditional Variable: Signal is stateless and has no record...

# Case Study (Soln) Add some record?

```
1 int* a, record = 0; // 1 -> finished, 0 -> not yet
2 pthread_cond_t cond;
3 pthread_mutex_t lock;
4
5 void* thr_func(void* arg) {
6     for (int i = 0; i < 100; i++) { a[i] = a[i] * 2; }
7     pthread_mutex_lock(&lock);
8     record = 1; // inform
9     pthread_cond_signal(&cond);
10    pthread_mutex_unlock(&lock);
11    sleep(3); // Do customization here
12    pthread_exit(NULL);
13 }
14
15 int main() {
16     a = malloc(100 * sizeof(int));
17     for (int i = 0; i < 100; i++) { a[i] = i; }
18
19     pthread_cond_init(&cond, NULL);
20     pthread_mutex_init(&lock, NULL);
21     pthread_t thr;
22     pthread_create(&thr, NULL, thr_func, NULL);
23
24     pthread_mutex_lock(&lock);
25     while (record != 1) { // not yet finished
26         pthread_cond_wait(&cond, &lock);
27     }
28     printf("All threads finished job and still alive\n");
29
30     pthread_join(thr, NULL);
```

## If Wait Before Post?

1. record = 0 != 1 so Main Thread call cond\_wait(&cond), put itself in waiting queue, being blocked and release lock
2. Child Thread acquire lock, update record = 1 and call cond\_signal(&cond) → wake up main and release lock
3. Main Thread being waked up and found record = 1 so leave loop → all done 🤔

Question: Will single mutex\_lock helps?

## Solution

Use a variable to record cond\_signal operation. (e.g., **an integer**)

This variable is shared among threads

→ Need **Mutex Lock** to protect thread safety:

- Line 7-10, when cond\_signal, i.e., update, also change the state. Both signal and change shall be protected by Mutex Lock.
- Line 24-27, when cond\_wait, i.e., receiving update, add **while (record != 1)** to check record. Shall be protected by Mutex Lock → **pthread\_cond\_wait has Mutex Lock as the 2nd parameter**

## If Post Before Wait?

1. Child Thread acquire lock, update record = 1, call cond\_signal(&cond) → found queue empty and release lock
2. Main Thread acquire lock and found record = 1 so leave loop → all done 🤔

Problem Solved!

But Cond Variable + (Integer) State = \_\_\_\_\_?

# Other Semaphore APIs

Get the Semaphore Value

```
int sem_getvalue(sem_t* sem, int* sval)
```

Place the current value of *sem* to *sval*.

```
int sem_trywait(sem_t* sem)
```

The unblocked version of `sem_wait`, will:

- return 0 if sem -= 1 succeed, i.e., sem > 0
- return non-zero if failed to do so

Named Semaphore (For Your Info)

For safety, macOS only support named semaphore.

```
sem_t* sem_open(const char* name, int oflag,  
mode_t mode, int value)
```

Open named semaphore by specifying name, mode, named semaphore is used for multi-processing

- `int sem_close(sem_t* sem)`
- `int sem_unlink(sem_t * sem)`

Note: Not recommending using semaphore in macOS → Underlying implementation is different



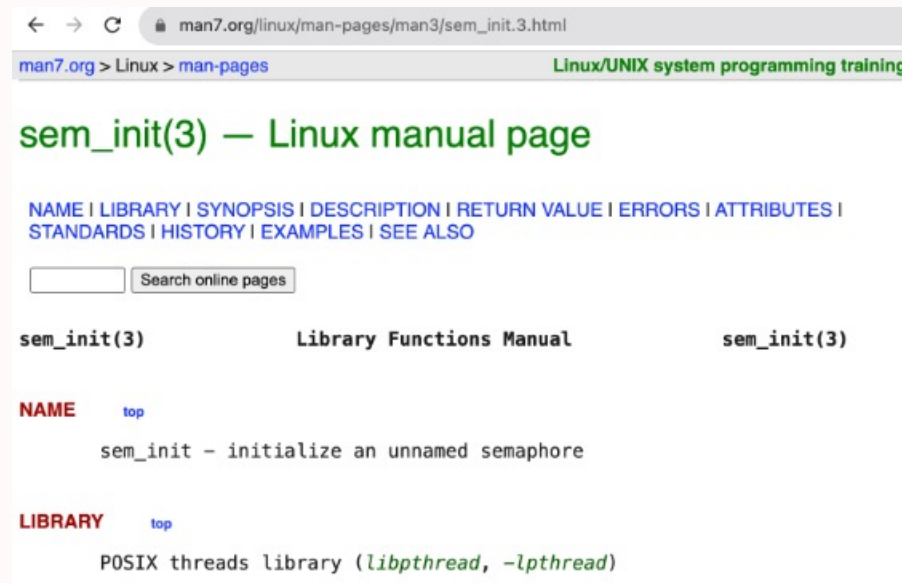
# Wanna know more?

In Linux, all these underlying API is provided by Linux Kernel and GNU C Library (glibc)

Search API Documentations

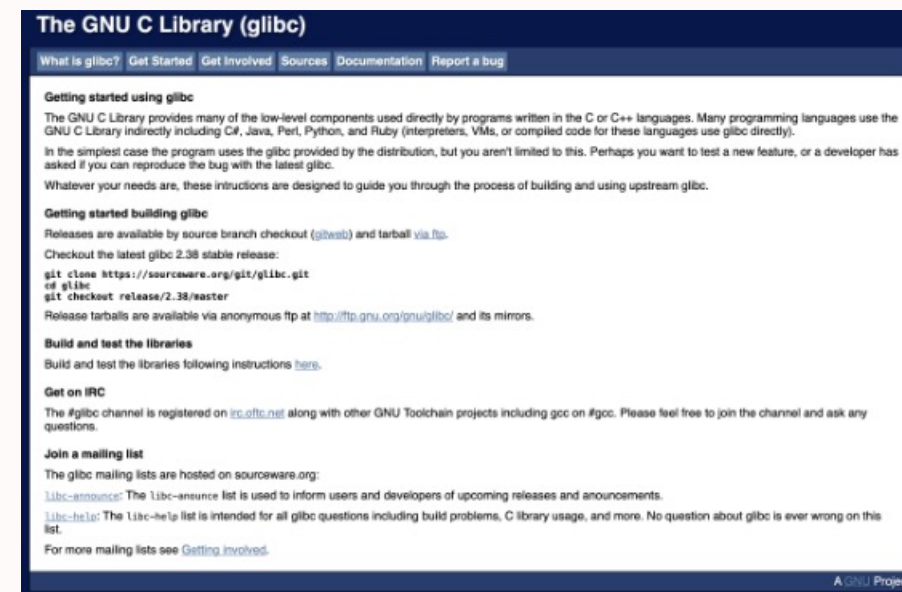
Recommended way → Linux Manual Page

Search **man [API]**, Go to link with **man7.org**



Search the source code

For corner cases, if don't violates the Principle of Operating System, no API record. Depends on actual implementation in [Linux](#) / [glibc](#):



Can also check C Language APIs

# Take Away

- Semaphore is widely used for Synchronization
  - Creation & Destroy: `sem_t, sem_init, sem_destroy`
  - Post & Wait: `sem_post, sem_wait`
  - Semaphore has an **atomic-protected integer state** as an record
- Conditional Variable is stateless → signal is just one-pass
  - Standalone Cond Variable **succeed Wait→Signal** but **fail Signal→Wait**
  - **Adding shared variables** helps → **Always used with mutex\_lock**
- Think Questions — Previously we discussed one-to-one sync, and how about:
  - One-to-many Sync?
  - Many-to-One Sync? → (while instead of if)
  - Many-to-Many Sync? → (while instead of if)

# Debugging Multi-Threading Program

- VS Code Setup
- Sample Debugging